

Mastering Kubernetes Taints & Tolerations



A Platform Engineer's
Complete Guide

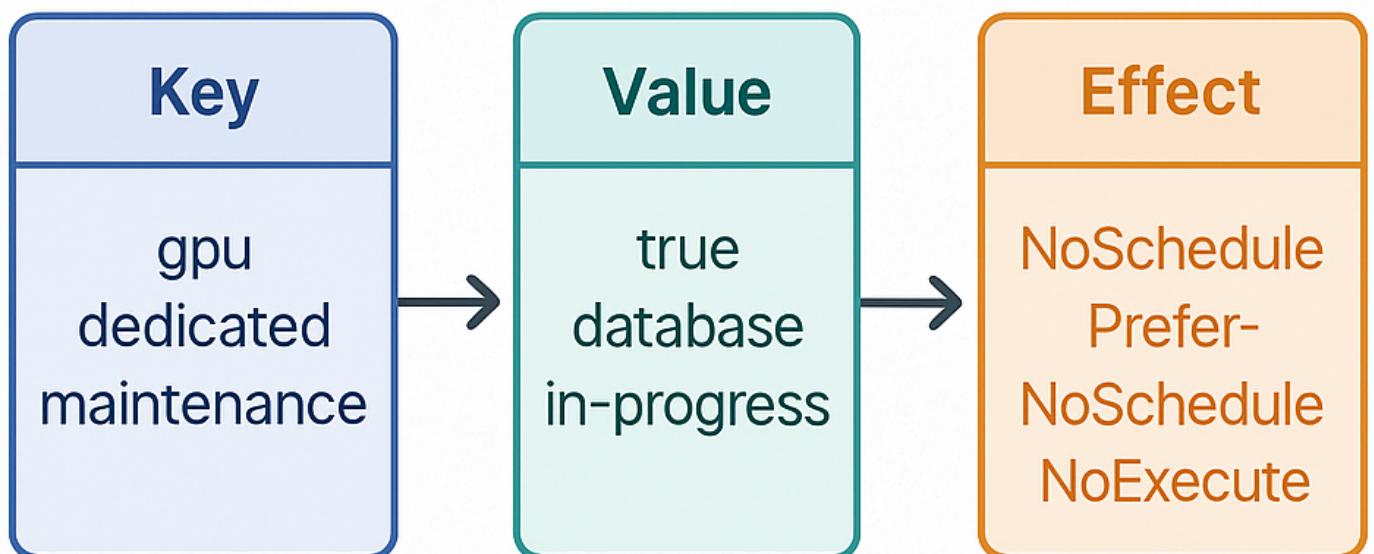
Swipe for insights →

What Are Taints?



-  Repel pods from specific nodes
-  Label-like properties applied to nodes
-  Signal to scheduler about node restrictions

Taint Anatomy 🔎



Taint Effects Explained



NoSchedule

Won't schedule new pods
(existing stay)



PreferNoSchedule

Tries to avoid scheduling
(soft rule)

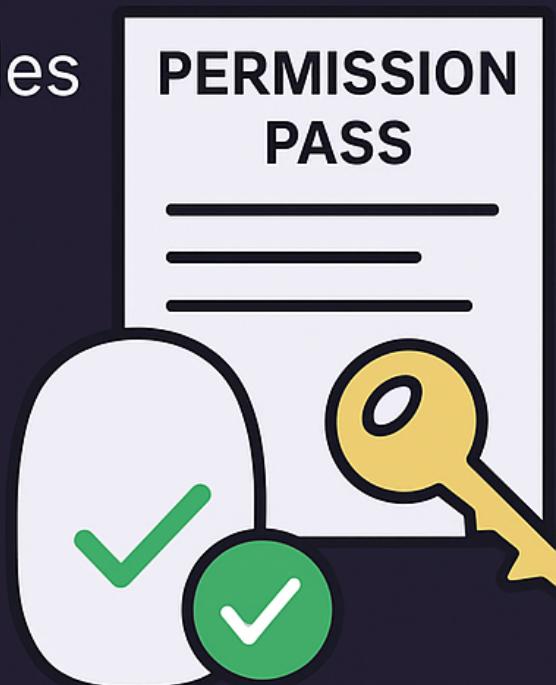


NoExecute

Evicts existing + blocks
new pods

What Are Tolerations?

- Allow pods on tainted nodes
- Defined in PodSpec
- Permission slip to bypass taints
- Must match taint properties



Toleration Structure



```
yamltolerations:
```

```
- key: "gpu"
```

```
operator: "Equal"
```

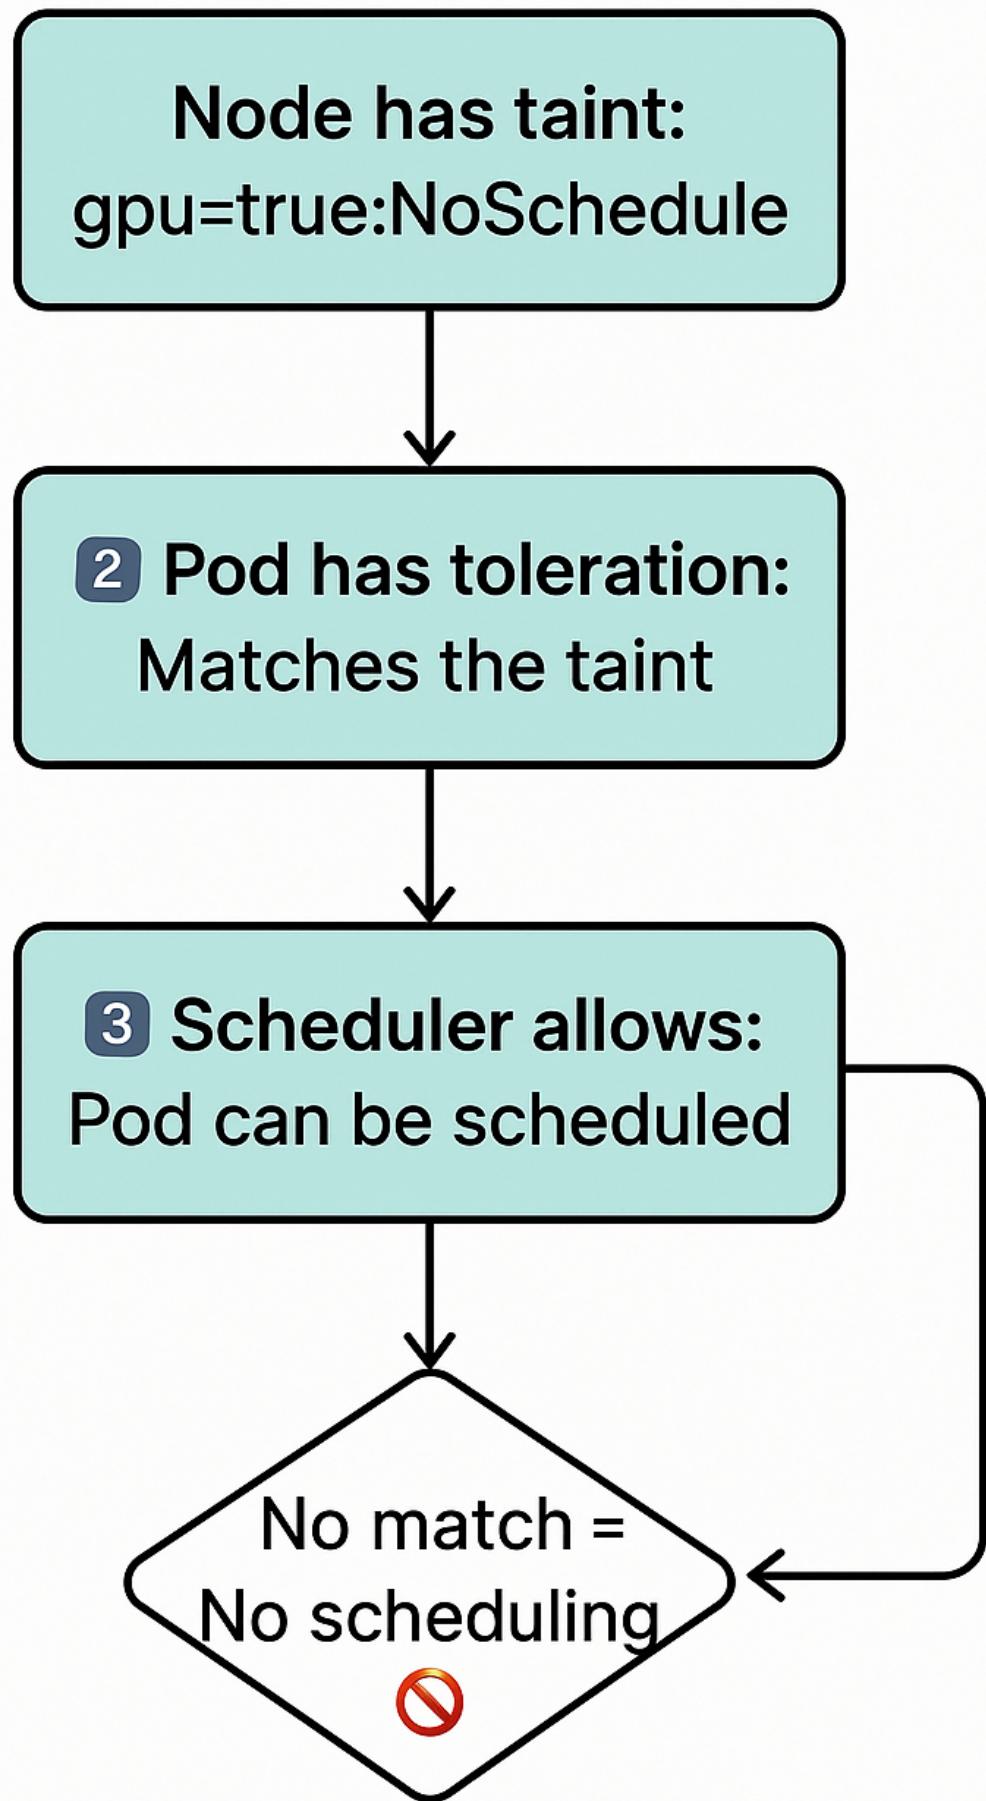
```
# or "Exists"
```

```
value: "true"
```

```
effect: "NoSchedule"
```

```
tolerationSeconds: 300
```

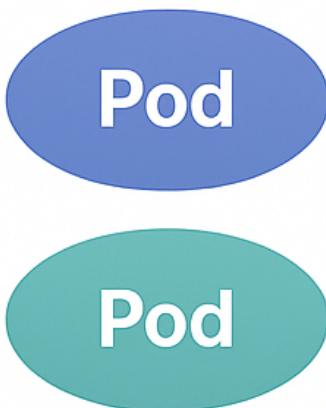
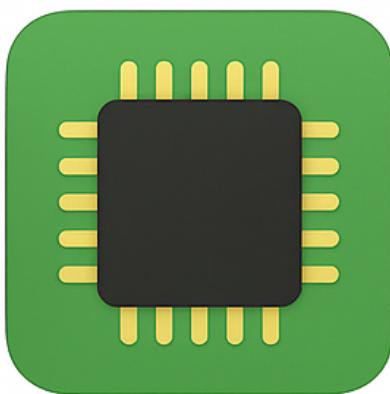
The Magic Happens Here ✨



Dedicate GPU Nodes



BEFORE



AFTER



Problem:
GPU nodes used by regular pods

Solution:
Taint GPU nodes:
`gpu=true:NoSchedule`
Add tolerations to
ML/AI pods

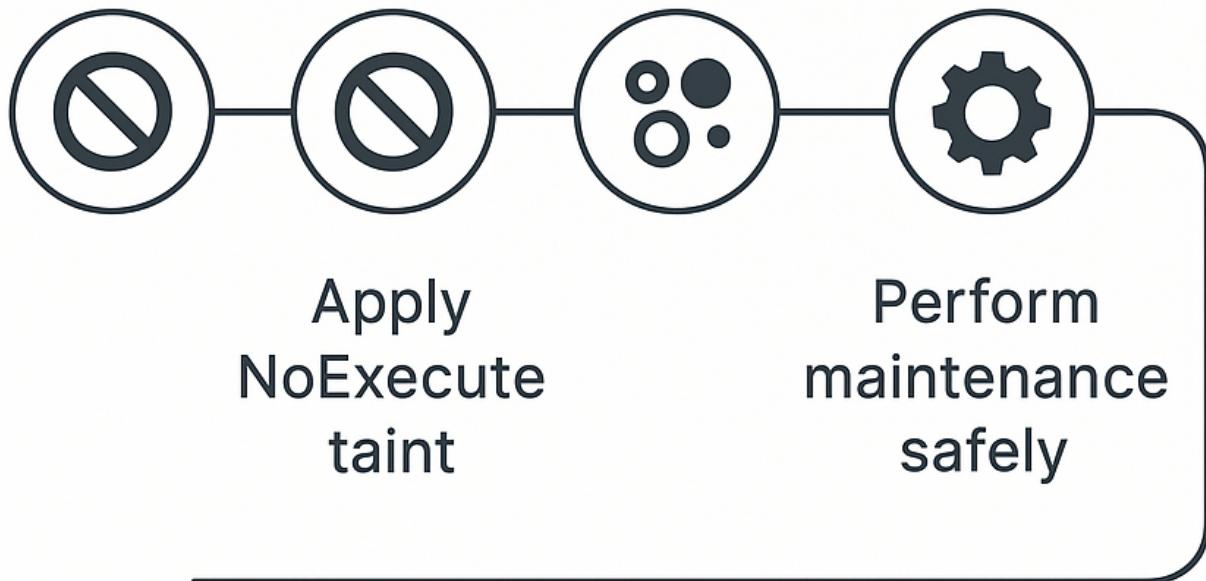
Result:
GPU reserved for intended workloads

Graceful Node Maintenance



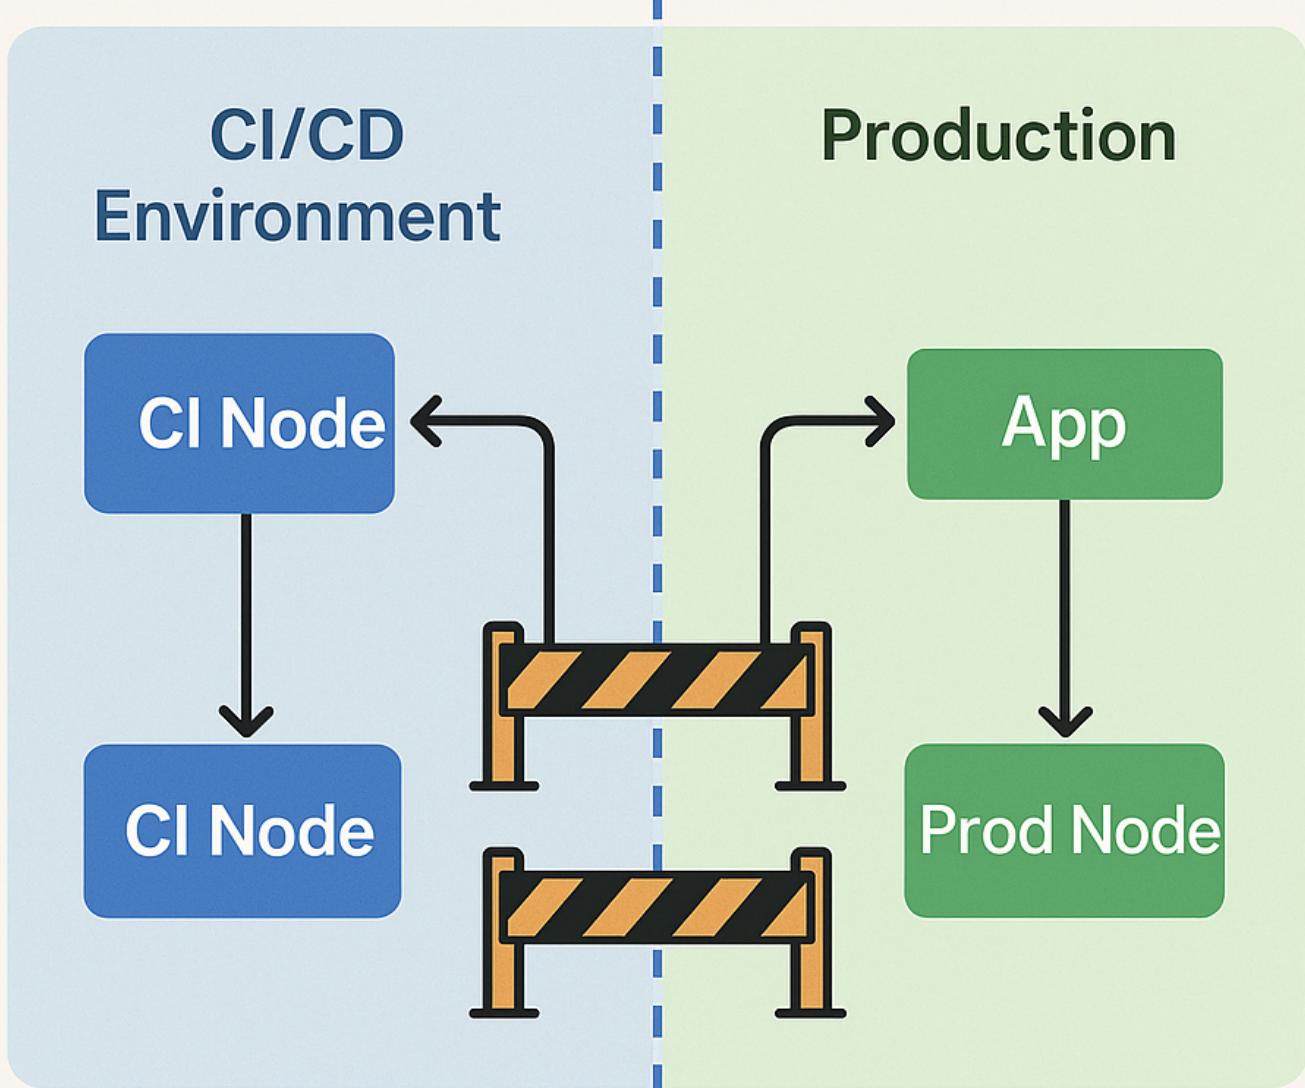
Cordon
the node

Pods with
tolerations get
tolerationSeconds



Graceful
eviction
happens

Isolate CI/CD Workloads



Why: Unpredictable resource usage

How: Taint CI nodes: `workload=cicd`:
`NoSchedule`

CI runners have matching tolerations

Benefit: Production workloads stay safe



Essential Commands

• Apply taint

```
kubectl taint nodes node-01  
key=value:Effect
```

• Remove taint

```
kubectl taint nodes node-01  
key=value:Effect-
```

• View taints

```
kubectl describe node node-01  
| grep Taints
```

NoExecute + tolerationSeconds



When NoExecute taint is added:

Pods without tolerations → Immediate eviction

Pods with tolerations → Stay indefinitely

With tolerationSeconds → Graceful period



Immediate
eviction



Stay
indefinitely



Graceful
period

Control Plane Protection



Default taint:

`node-role.kubernetes.io/control-plane:NoSchedule`

Purpose: Reserve resources for cluster management

Exception: System pods have tolerations

Best Practices

-  Clear naming conventions
-  NoSchedule for dedication
-  NoExecute for eviction
-  Combine with node affinity
-  Manage as code (GitOps)
-  Regular auditing

Avoid These Mistakes!



- ✖ Broad tolerations everywhere
- ✖ Forgetting tolerationSeconds
- ✖ Mismatched key/value pairs
- ✖ No documentation
- ✖ Testing only in production

Troubleshooting 101



Pod stuck in Pending?

Check node taints

Verify pod tolerations

Review scheduler events

Confirm resource availability

Advanced Patterns



- ◆ Empty key + Exists operator
- ◆ Empty effect matching
- ◆ Multiple taint combinations
- ◆ Custom node conditions
- ◆ Webhook-based tainting

Production Example



```
# GPU Node Taint
```

```
kubectl taint nodes gpu-node-01  
nvidia.com/gpu=tesla-v100:  
NoSchedule
```

```
# ML Pod Toleration
```

```
tolerations:  
key: "nvidia-com/gpu"  
value: "tesla-v100"  
effect: "NoSchedule"
```

Key Takeaways



- 1 Taints repel, tolerations allow**
- 2 Three effects: NoSchedule,
PreferNoSchedule, NoExecute**
- 3 Perfect for: Dedication,
maintenance, isolation**
- 4 Combine with other scheduling
features**
- 5 Test thoroughly before
production**

Want to Learn More?



Full guide with YAML examples:

<https://medium.com/@salwan.mohamed/a-comprehensive-guide-to-kubernetes-taints-and-tolerations-58af8659e92a>



Questions?
Drop them below!

Comment



Repost if this
helped you

Repost



Tag a colleague
who needs this

Tag

@SalwanMohamed