




MASTERING KUBERNETES

MULTI-CONTAINER DESIGN PATTERNS

A Complete Guide
for Platform Engineers

Build Scalable,
Maintainable Container
Architectures




THE PROBLEM



MOST DEVELOPERS THINK:

Pod = One Container

But this limits Kubernetes' true power!



- ✖ Single container doing everything
- ✖ Tight coupling
- ✖ Hard to maintain
- ✖ Limited reusability




THE SOLUTION

✓ THE REALITY:

Pods are Shared Execution Environments

Multiple containers working together
as a cohesive unit

- Separation of concerns
- Enhanced modularity
- Reusable components
- Unix philosophy:
“Do one thing well”





WHAT IS A POD?


Shared Network

- Same IP address & port space




Shared Storage

- Volumes & filesystem communication



Shared Lifecycle

- Start, stop, scale together






Pattern Overview


1 3 ESSENTIAL MULTI-CONTAINER PATTERNS

SIDECAR PATTERN




Helper container
extends
functionality

AMBASSADOR PATTERN



Proxy container
handles connections

3 ADAPTER PATTERN



Transform container
converts output


Each solves specific architectural challenges

SIDECAR PATTERN

Deep Dive

SIDECAR PATTERN: The Swiss Army Knife

Helper container alongside
main application



COMMON USE CASES:

-  Logging agents (Fluentd, Filebeat)
-  Monitoring agents (Prometheus)
-  Security proxies (Istio)
-  Config management
-  Secret rotation



SIDECAR IN ACTION

volumes:

 name: shared-logs

 emptyDir: ()

Shared: Volume
for communication

Both containers
mount the same
volume!





AMBASSADOR PATTERN

The Proxy Layer




PERFECT FOR:

- Service discovery abstraction
- Connection pooling
- Load balancing
- Protocol translation
- Circuit breaking & retry logic

ADAPTER PATTERN

The Translation Layer



Container transforms output
for external systems

USE CASES:

- Log format transformation
- Metrics format conversion
- API response adaptation
- Protocol conversion
- Data normalization

LET'S BUILD A LOGGING SIDECAR




ARCHITECTURE:

1. Bash log generator (main)
2. Log processor (sidecar)
3. Shared emptyDir volume



Real implementation
with Docker & K8s!



docker®





MAIN APPLICATION: Log Generator



```
1 #!/bin/bash
2 while true; do
3     MESSAGE= "$(date -Iseconds) -
4         GUID: $(uuidgen)"
5     echo $(MESSAGE) >> /app/logs/
6         my-test-log.log
7     sleep $(shuf -i 1-3 -n 1)
8     done
9 done
```

Generates timestamped logs with UUIDs



SIDECAR: Log Processor

```
#!/bin/bash
tail -f /var/log/myapp/*.log
i while read line; do
if [[ $line == [a-f0-9-]{36}]]:
then
$(BASH_REMATCH[0])
fi
done
```

[a-f0-9-]{36}

123e4567-e89b-12d3-a456-4266544000

Extracts UUIDs from log
stream in real-time



POD CONFIGURATION



```
apiVersion: v1
kind: Pod
metadata:
  name: test-logger
spec:
  volumes:
    name: shared-logs
  emptyDir: {}
  volumeMounts
    name: shared-logs
  mountPath: /app/logs
```

RESOURCE CONFIGURATION

RESOURCE ALLOCATION



containers:


my-test-logger

resources:

 memory: 200Mi 

 cpu: 600m

- log-reader

 memory: 200Mi 

 cpu: 600m

Proper resource boundaries
prevent conflicts



🚀 DEPLOYMENT & VERIFICATION

```
# Deploy the pod  
kubectl apply -f / pod.yaml  
  
# Check status  
kubectl describe pod test-logger  
  
# View sidecar logs  
kubectl logs pod/test-logger  
-c log-reader --follow  
  
✓ Live log streaming from  
the sidecar!
```



INIT CONTAINERS

The Setup Crew

Run **BEFORE** main containers start



**INIT
CONTAINER**

PERFECT FOR:

- Database migrations
- Configuration setup
- Dependency verification
- Security scanning
- Environment preparation

.....>



**MAIN
CONTAINER**

initContainers run sequentially, then
main containers start


SERVICE MESH WITH SIDECAR



Istio automatically injects proxy sidecars

containers:

- name: my-app
image: my-app:latest
- name: istio-proxy
image: istio/proxy2:1.20.0



Automatic traffic management,
security, observability

COMPREHENSIVE MONITORING

Monitoring Stack



web-app

my-web-app:latest



metrics-exporter

prometheus/nd-exporter



log-shipper

fluent/fluent-bit



trace-collector

jaegertracing/jaeger agent



Full observability in one pod!

BEST PRACTICE



✓ DESIGN PRINCIPLES


1 Single Responsibility

Each container =
one clear purpose




2 Resource Boundaries

Define appropriate limits



3 Failure Isolation

Containers fail
independently



4 Security First

Non-root users,
minimal images
network policies



COMMON PITFALLS



AVOID THESE MISTAKES

- ✗ Tight coupling between containers
- ✗ Resource competition
- ✗ Complex log management
- ✗ Poor debugging strategies



SOLUTIONS:



Well-defined interfaces



Proper resource allocation



Standardized logging



Comprehensive monitoring



PERFORMANCE OPTIMIZATION

PERFORMANCE CONSIDERATIONS



NETWORK

- Localhost communication (very fast)
- No encryption overhead
- Shared network namespace



STORAGE

- emptyDir for fast ephemeral storage
- tmpfs for high-performance temp storage



RESOURCES

- CPU sharing strategies
- Memory allocation patterns



TESTING STRATEGIES


TESTING MULTI-CONTAINER PODS

UNIT TESTING

>_

```
docker run -rm  
my-app:test  
npm test
```

INTEGRATION TESTING




```
kubectl apply -f test-pod.yaml  
kubectl wait --for=condition=ready  
pod/test-pod
```

LOAD TESTING



Test combined resource usage under load



Each combiner
AND
their inter-
actions!



Each container
AND their interactions!

MIGRATION ROADMAP

FROM SINGLE TO MULTI-CONTAINER:



FROM

TO

- 1 Identify separable concerns
- 2 Extract functionality
- 3 Define communication interfaces
- 4 Gradual rollout with feature flags
- 5 Monitor and validate

Start small, think big!



WHAT'S NEXT?



WEBASSEMBLY (WASM):

Lightweight alternatives
to sidecars



SERVERLESS MULTI-CONTAINER:

Knative exploring pod-level
scaling



EDGE COMPUTING:

Multi-container patterns
at the edge

The future is multi-container
native!



KEY TAKEAWAYS

- ✓ Think in terms of PODS, not just containers
- ✓ Embrace SEPARATION OF CONCERNS
- 🚀 Design clear COMMUNICATION interfaces
- ✓ Monitor at CONTAINER & POD levels
- ✓ Start small, scale thoughtfully

Ready to transform
your architecture?



Like if this
helped!



Share with
your team!



Comment
your
experiences!

#Kubernetes #DevOps #PlatformEngineering