



ArgoCon

EUROPE



One Engine to Rule Them All

Unifying Cloud Workloads with Argo Workflows

Sebastian Beyvers, Giessen University



ArgoCon

EUROPE

1 April 2025

London, England



**Sebastian
Beyvers**
Distributed Systems
Researcher
Giessen University

Agenda



Introduction

Background: where are we coming from



Our Problem

A heterogenous ecosystem of workflow engines and compute tools

Solution

Argo workflows as universal scheduling engine ?

Implementation

Migration and Integration patterns

Results

What are the benefits ? What are the drawbacks ?

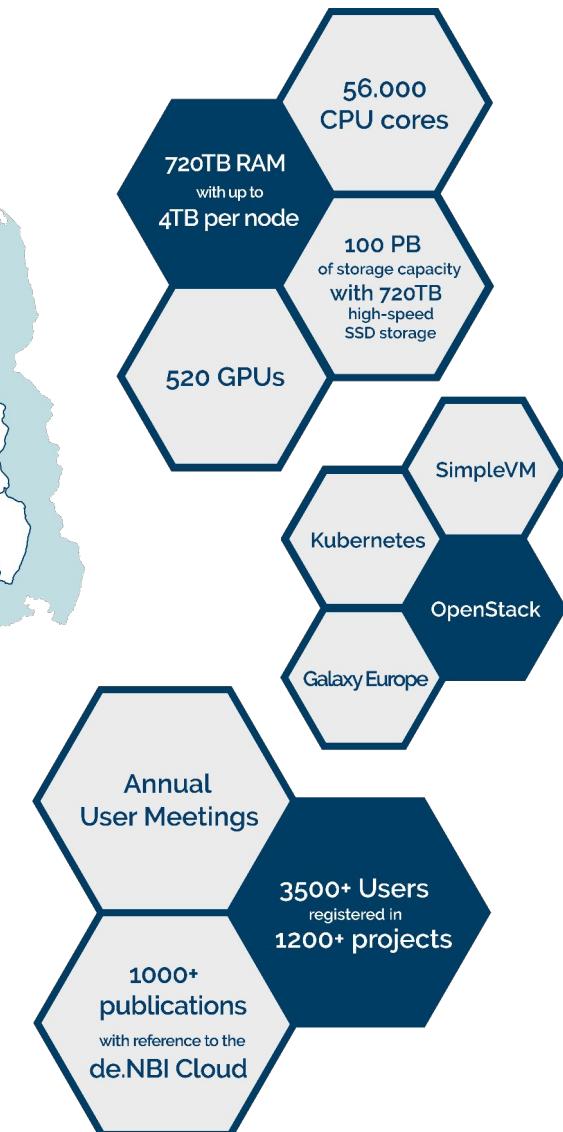
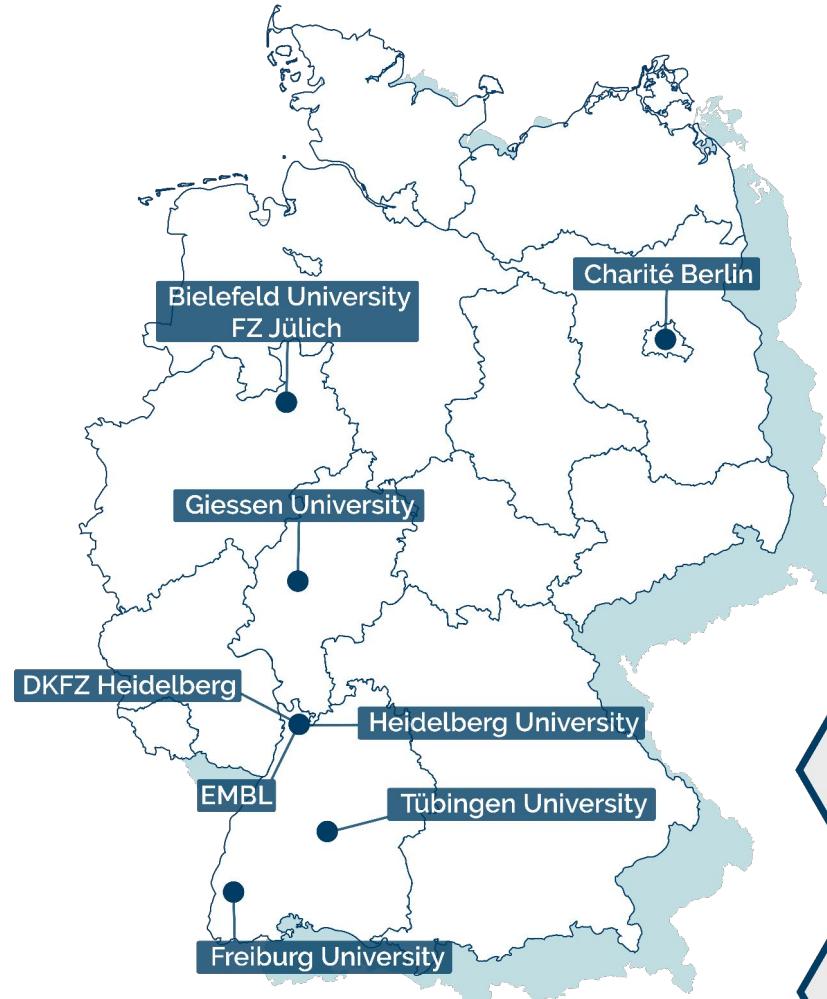
Background



GERMAN NETWORK FOR BIOINFORMATICS INFRASTRUCTURE



- Largest community cloud for life-science in Germany
- Free to use for the research community
- OpenStack
- Managed Kubernetes
- SLURM
- VMs

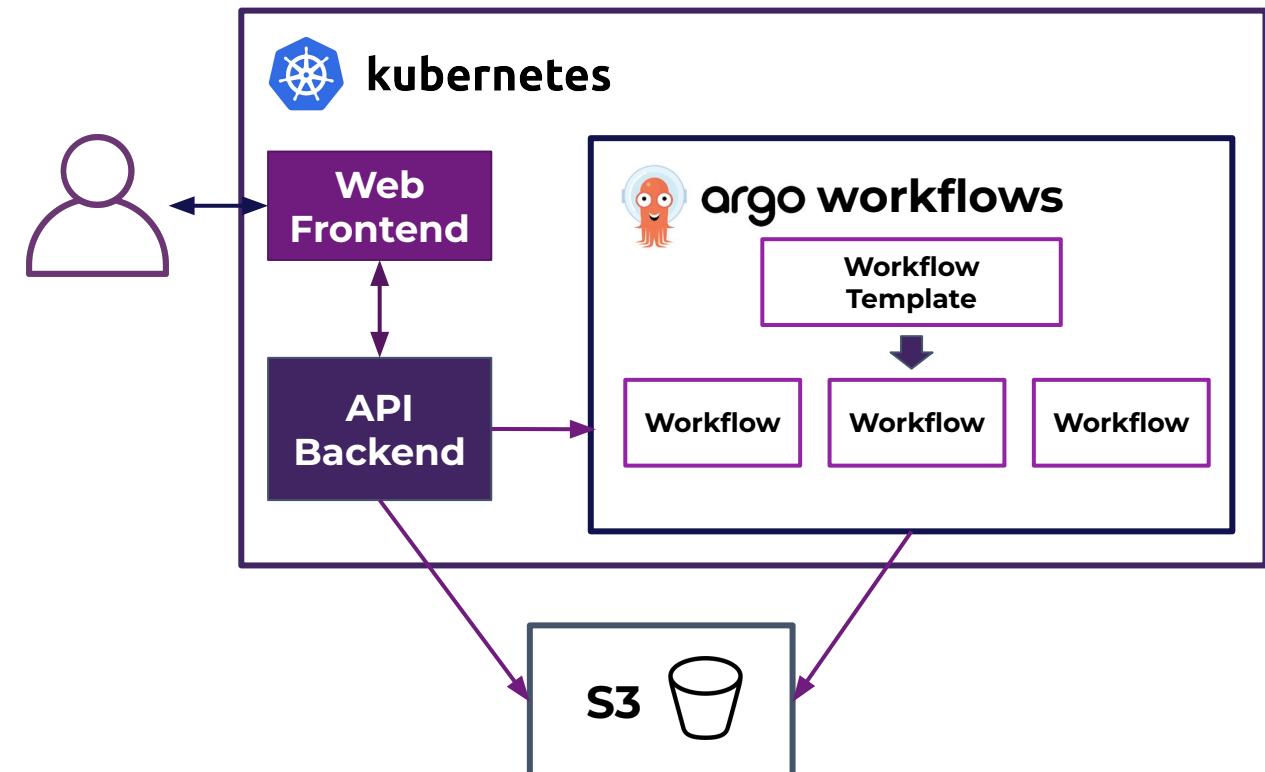


Our Argo Workflows journey



- Started ~5 years ago
- Currently ~10 clusters:
 - 200 - 5000 vCPUs
 - Including some high-mem nodes > 2 TB RAM
- High throughput with 100-1000s workflows per day
- Long running workflows (Days - Weeks)
- Primarily analysis workflows, almost no streaming / real time analysis

	NAME	NAMESPACE	STARTED	FINISHED	DURATION	PROGRESS
☐	edgar-job-	bakta	4d7h ago	1d6h ago	3d1h	17632/1...
☐	bbb7465a-6cd1-4924-b40d-333e2417af0c-cdlt4					



Why Argo Workflows ?



ArgoCon
EUROPE

Flexibility

You can do almost everything with Argo workflows

Extensibility

Plugin system, Workflow Templates and easy API access allows for scheduling of pre-configured workflows

Events

Argo Events provides additional automation capabilities and integration with message queues like Kafka, RabbitMQ etc.

Kubernetes native

A K8s native system enables easy handling of Kubernetes resources and optimal use of containers

Ecosystem

The Argo Ecosystem provides many additional useful tools like the **Hera** python library, integration with **Argo CD**, delivery with **Argo Rollouts**

The Problem



Cloud

Argo
workflows

Nextflow

Snakemake

KubeFlow

Apache
Airflow

Spark

Custom
Jobs

HPC / SLURM

Snakemake

CWL

Spark

Conveyor

Nextflow

Custom
Jobs

VMs

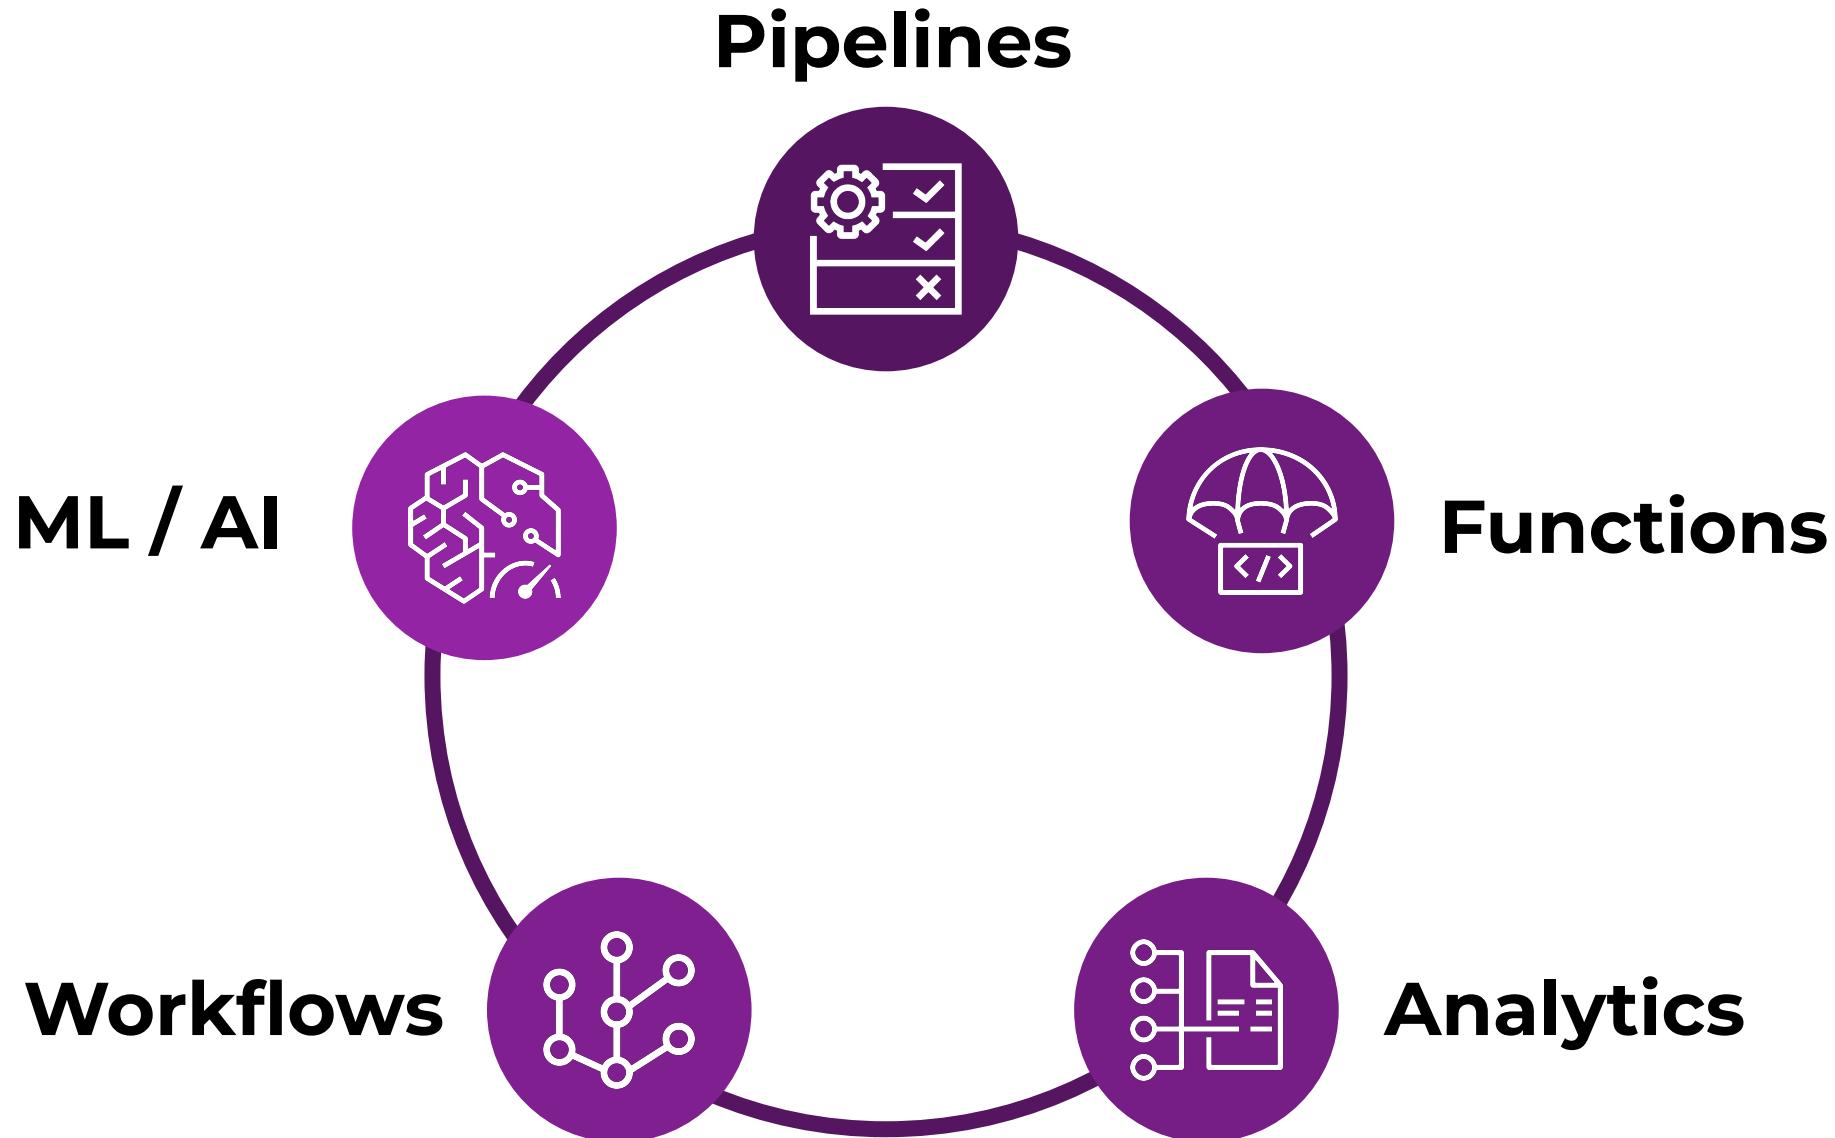
Snakemake

Custom
Jobs

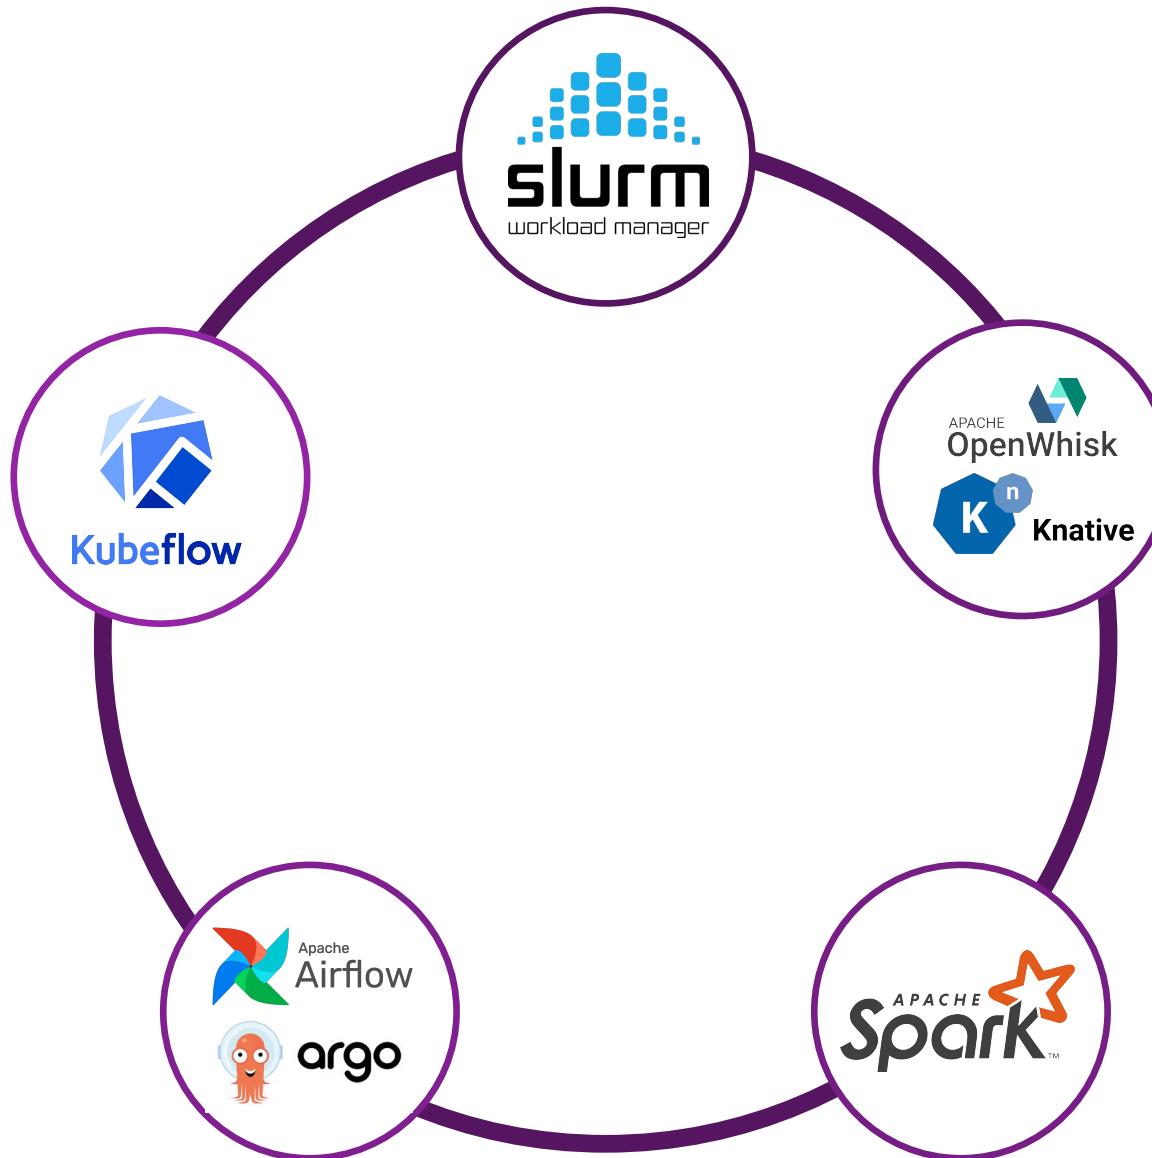
Nextflow

A wide range of workflow tools and environments

Promises of a public cloud



DIY: Collection of tools



DIY: Pros & Cons



ArgoCon
EUROPE



Pro

- Sovereignty & Independence
- Specialized solution
- Data protection
- Flexibility and cost effectiveness

Con

- Mental overhead
- Larger human resource requirements
- Drastically increased onboarding time



Argo Workflow as coordinator



**Can Argo workflows be a universal
scheduling engine ?**

**How can Argo Workflows
coordinate my cloud workloads ?**

How can Argo Workflows coordinate my cloud workloads ?

A: Migration



How can Argo Workflows coordinate my cloud workloads ?

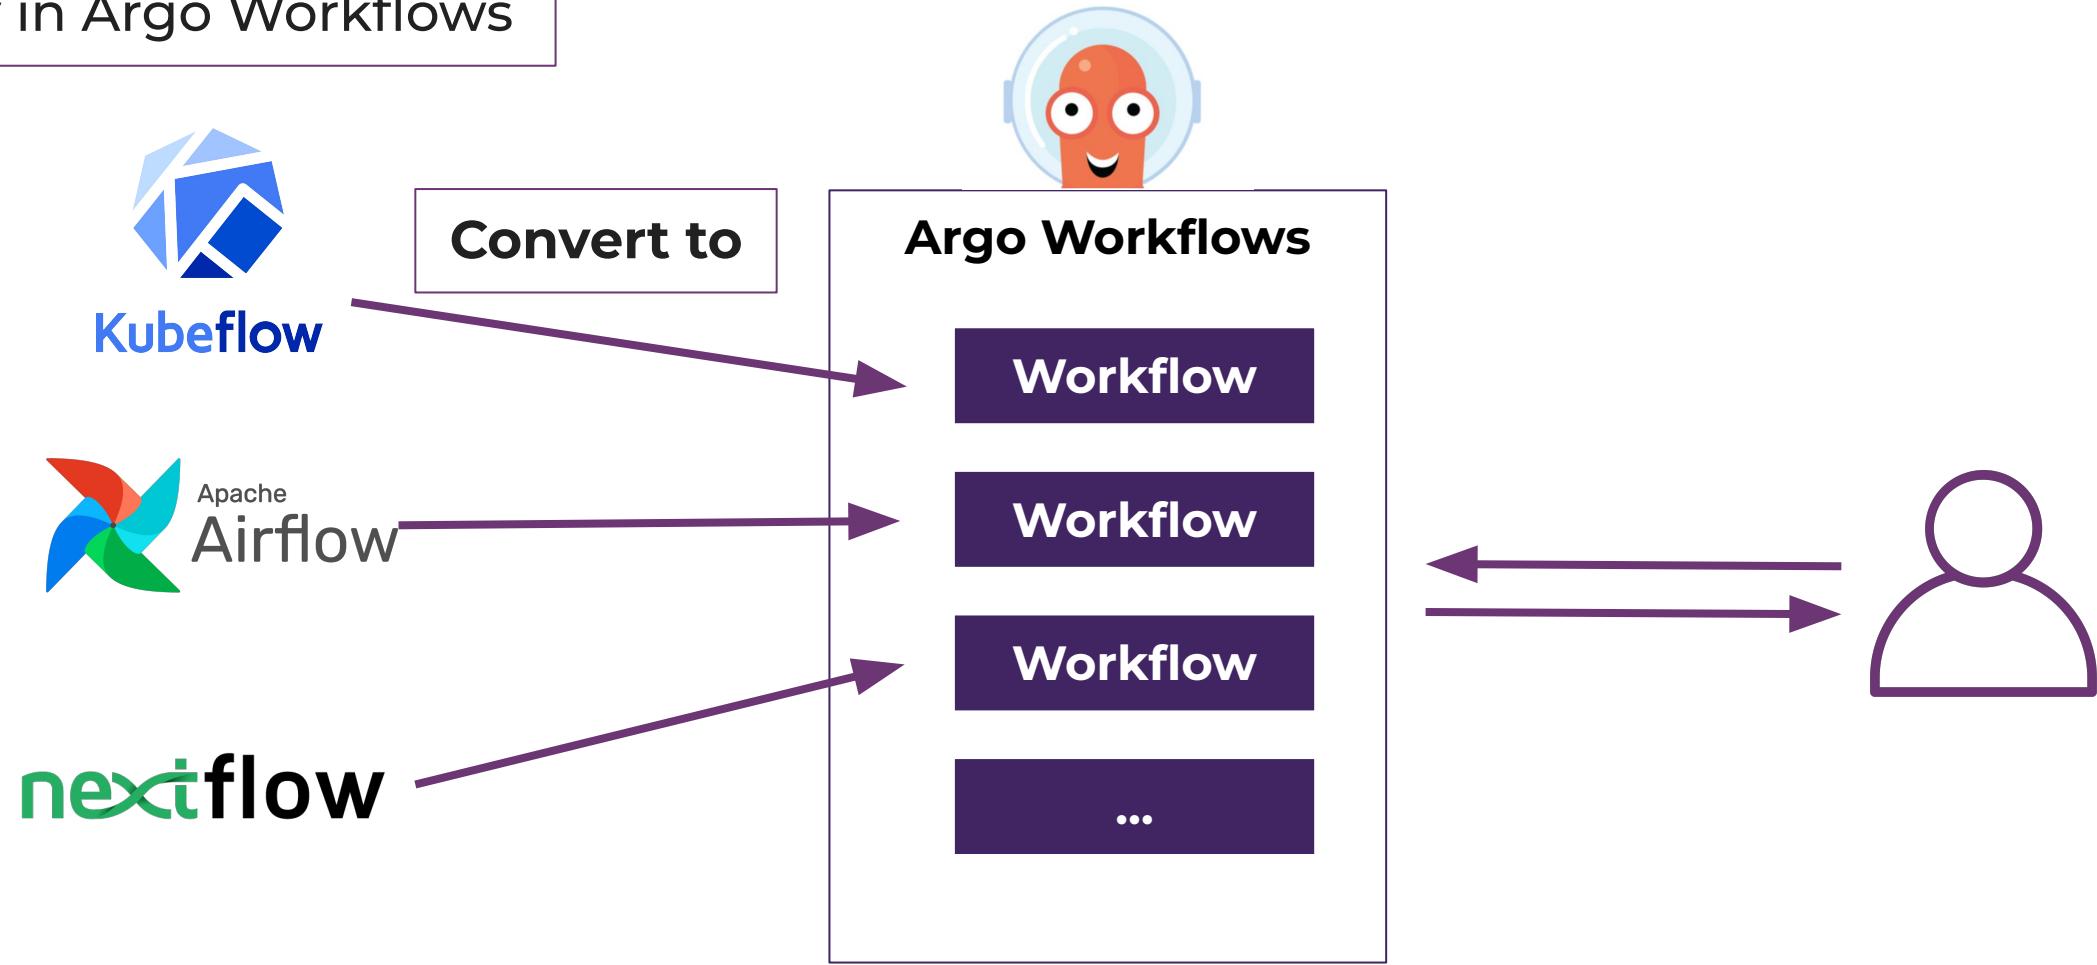
A: Migration

B: Integration

Strategy: Migration



Reimplement the workload directly in Argo Workflows



Strategy: Integration



ArgoCon
EUROPE

Wrap the existing system
in an Argo workflow



Init

Main

Finalize

- | | | |
|--|--|---|
| <ul style="list-style-type: none">• Initializes data storage• Upload data to a shared storage• Pre-configure variables and secrets | <ul style="list-style-type: none">• Start the external run• Monitor the status• Mirror the external status in the argo workflow status | <ul style="list-style-type: none">• Accumulate results• Process outcome• (optional) Upload results to cloud storage |
|--|--|---|

Strategy: Migration vs. Integration

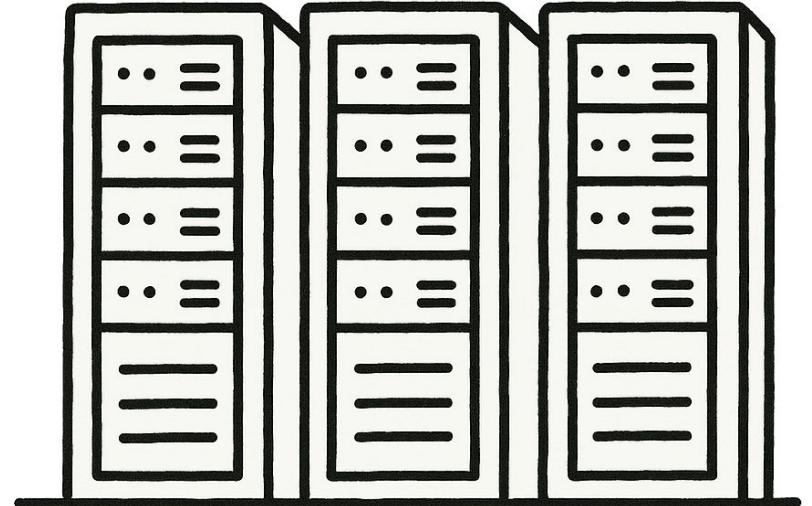


		Integration	Migration
Hard requirement to use an existing workflow engine?	yes	✓	✗
Is the existing workflow very complex ?	no	✗	✓
Sufficient team expertise for the existing solution ?	yes	✓	✗
Has the workload a varying degree of resource requirements ?	no	✗	✓
Is your workload already containerized or easily containerizable ?	yes	✗	✓
	no	✓	✗

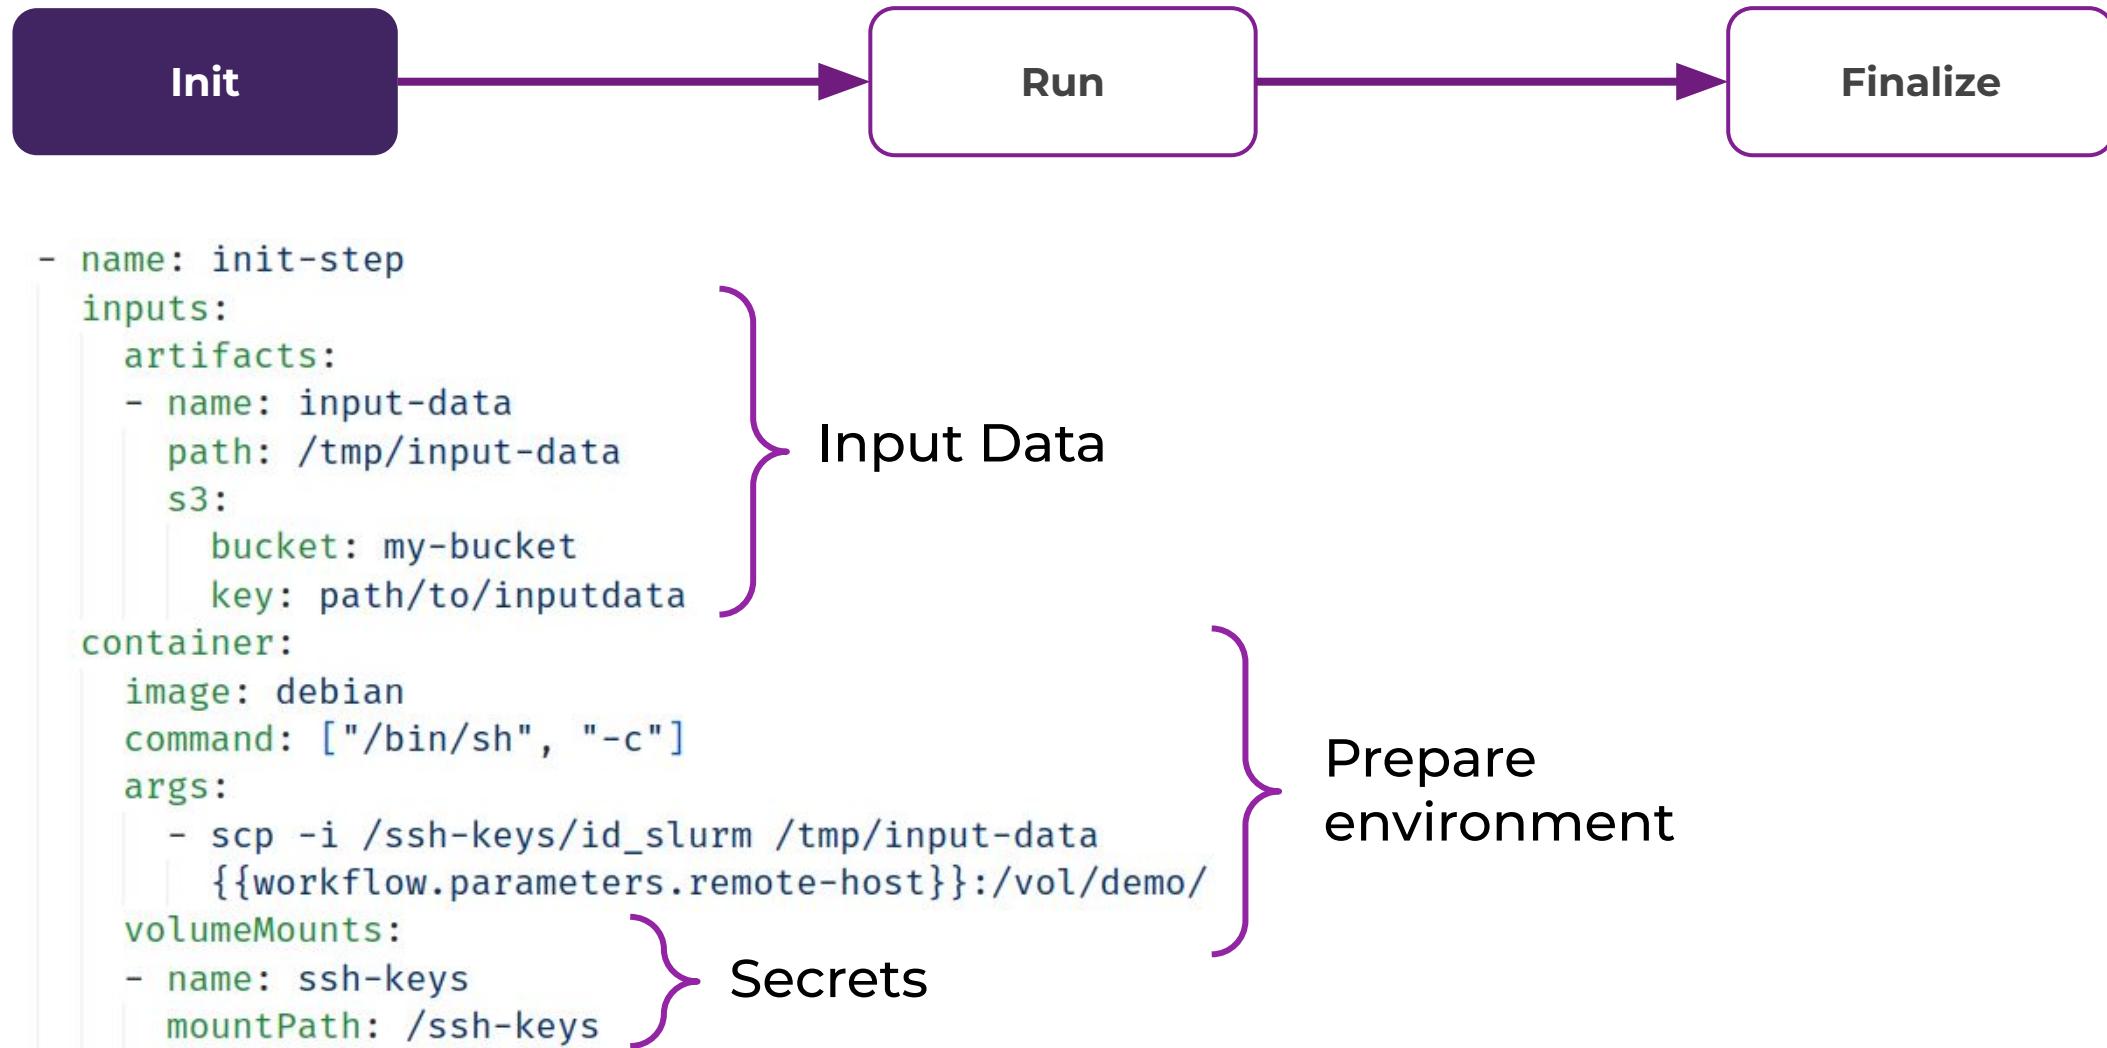
HPC Workloads: Migration



- Rewrite the analysis workflow as Argo workflow
- **Challenges:**
 - Shared file system
 - Shared permissions
 - Internal databases
 - Specialized Hardware (FPGA, GPUs etc.)



HPC Workloads: Integration



HPC Workloads: Integration



```
- name: run-step
  container:
    image: debian
    command: ["/bin/sh", "-c"]
    args:
      - ssh -i /ssh-keys/id_slurm
        {{workflow.parameters.remote-host}}
        'sbatch --wait /vol/demo/job.sh'
    volumeMounts:
      - name: ssh-keys
        mountPath: /ssh-keys
```

Run SLURM workload in HPC environment

HPC Workloads: Integration



```
- name: finalize-step
  container:
    image: debian
    command: ["/bin/sh", "-c"]
    args:
      - scp -i /ssh-keys/id_rsa
        {{workflow.parameters.remote-host}}:/vol/demo/results
        /tmp/results
    volumeMounts:
      - name: ssh-keys
        mountPath: /ssh-keys
    outputs:
      artifacts:
        - name: results
          path: /tmp/results
          s3:
            bucket: my-bucket
            key: path/to/results
```

Retrieve remote data from HPC environment

Create output artifact from remote data

Data Pipelines: Migration



ArgoCon
EUROPE



```
1 from airflow.decorators import dag, task
2
3 @task()
4 def hello(name: str):
5     print(f"Hello {name}!")
6
7 @dag(
8     dag_id = "dag_example",
9 )
10 def dag_example():
11     t1 = hello(name="hello-1")
12     t2 = hello(name="hello-2")
13     t1 >> t2
14
15 dag_instance = dag_example()
```

```
1 from hera.workflows import DAG, Parameter, Workflow, script
2
3 @script()
4 def hello(name: str):
5     print(f"Hello {name}!")
6
7 with Workflow(
8     generate_name = "dag-example",
9     entrypoint = "dag",
10 ) as w:
11     with DAG(name = "dag") as d:
12         t1 = hello(name="hello-1", arguments = {"name": "hello-1"})
13         t2 = hello(name="hello-2", arguments = {"name": "hello-2"})
14         t1 >> t2
```

Data Pipelines: Integration



```
- name: trigger-airflow
  script:
    image: python:3.14
    command: [python]
    source: |
      import requests, time, sys
```



```
airflow_url = "{{workflows.parameters.airflow_url}}"
dag_id = "{{workflows.parameters.dag_id}}"
run_id = "{{workflows.parameters.run_id}}"
headers = "{{workflows.parameters.headers}}"
payload = {"dag_run_id": run_id}
```

```
response = requests.post(f"{airflow_url}/api/v1/dags/{dag_id}/dagRuns",
                           json=payload, headers=headers)
if response.status_code != 200:
    print(f"Failed to trigger DAG: {response.status_code}")
    sys.exit(1)
```

```
status_endpoint = f"{airflow_url}/api/v1/dags/{dag_id}/dagRuns/{run_id}"
while True:
    status_response = requests.get(status_endpoint, headers=headers)
    state = status_response.json().get('state')
    if state in ['success', 'failed']:
        sys.exit(0 if state == 'success' else 1)
    time.sleep(5)
```

} Setup parameters

} Run DAG

} Wait for completion

Analytics: Integration



```
- name: sparkapp-operator
resource:
  action: create
  successCondition: status.applicationState.state == COMPLETED
  failureCondition: status.applicationState.state == FAILED
  manifest: |
    apiVersion: sparkoperator.k8s.io/v1beta2
    kind: SparkApplication
    metadata:
      name: spark-pi
    spec:
      type: Scala
      mode: cluster
      image: spark:3.5.1
      mainClass: org.apache.spark.examples.SparkPi
      mainApplicationFile: local:///opt/spark/examples/jars/spark-examples_2.12-3.5.1.jar
```

The code block is annotated with two curly braces. One brace groups the 'resource' section and its nested fields ('action', 'successCondition', 'failureCondition', 'manifest') under the heading 'Conditions'. Another brace groups the entire 'resource' section and the 'spec' field under the heading 'Kubernetes resource'.

- Using KubeFlow Spark operator:
 - <https://github.com/kubeflow/spark-operator>
- <https://pipekit.io/blog/argo-workflows-spark>

Functions: Migration



- Use different templates for programming languages
- Use Python, Javascript, Rust etc. for function
- Possible strategy: One **WorkflowTemplate** per programming language

```
- name: uv-example
  script:
    image: ghcr.io/astral-sh/uv:debian
    source: |
      #!/usr/bin/env -S uv run --script
      # /// script
      # requires-python = ">=3.11"
      # dependencies = [
      #   "httpx",
      #   "rich",
      # ]
      # ///

      import httpx
      from rich.pretty import pprint

      resp = httpx.get("https://peps.python.org/api/peps.json")
      data = resp.json()
      pprint([(k, v["title"]) for k, v in data.items()][:10])
```

Unified Developer Experience



- single entry point for all workflows
- centralized API
- KPI & metrics friendly
- Pre-configured workflows using **WorkflowTemplates**
- Fast onboarding time

“All under one roof”

	NAME	NAMESPACE	STARTED	FINISHED	DURATION	PROGRESS
<input type="checkbox"/>	Airflow feature-count-QkKBfNpGA6s2 [link]	argo	1h38m ago	1h12m ago	26m5s	3/3
<input type="checkbox"/>	SLURM nextflow-rna-seq-MT3ryeuzkrZ4	argo	2h42m ago	1hm44m ago	58m10s	3/3
<input type="checkbox"/>	SLURM snakemake-fastqc-RFB5RXHwe3bD	argo	3h10m ago	3h1m ago	9m2s	3/3
<input type="checkbox"/>	SLURM snakemake-fastqc-YawSCR57jpjp	argo	3h22m ago	3h5m ago	16m3s	3/3
<input type="checkbox"/>	Function: python python-tadrep-FjVq8wcJD4aT	argo	4h31m ago	4h22m ago	9m5s	3/3
<input type="checkbox"/>	SLURM nextflow-rna-seq-66qPJ2F3pKuC	argo	5h10m ago	4h59m ago	11m6s	3/3
<input type="checkbox"/>	Airflow feature-count-Q9GP5vdgtT2t [link]	argo	6h10m ago	5h40m ago	29m10s	3/3
<input type="checkbox"/>	Airflow feature-count-2Mc5QyAby7ek [link]	argo	6h44m ago	6h22m ago	21m10s	3/3

Limitations & Challenges



ArgoCon
EUROPE

Low latency

Stream processing

Performance overhead

Specialized hardware (GPUs etc.)

Multiple Argo Workflow instances

Argo workflows can be used as a unified scheduling engine

WorkflowTemplates can help to drastically reduce the onboarding time

Kubernetes **can replace** many traditional HPC workloads !

Outlook



Argo workflows wishlist:

- Better **synchronization** between different Argo workflow clusters
- Smarter resource **request / limit** configurability
- **Kueue** integration
- Better data management capabilities e.g. **FUSE**

Personal wishlist:

- Successful **migration** of most cloud workloads to Argo workflows
- Automated migration tools
- Integrate our data management solutions into workflows

Share your feedback



Join our poster session



KubeCon



CloudNativeCon
Europe 2025

Thursday April 3, 13:15 - 14:15 BST
@Level 1 N8-N9 | Poster Pavilion

Questions ? Contact me:

- ✉ sebastian.beyvers@cb.jlug.de
- 🐦 @St4NNi
- 💻 St4NNi

Thanks to our partners:

