

DevOps Helm Charts

June 24, 2025



DEVOPS HELM CHARTS

Copyright

All product technical documentation is
Ping Identity Corporation
1001 17th Street, Suite 100
Denver, CO 80202
U.S.A.

Refer to <https://docs.pingidentity.com> for the most current product documentation.

Trademark

Ping Identity, the Ping Identity logo, PingAccess, PingFederate, PingID, PingDirectory, PingDataGovernance, PingIntelligence, and PingOne are registered trademarks of Ping Identity Corporation ("Ping Identity"). All other trademarks or registered trademarks are the property of their respective owners.

Disclaimer

The information provided in Ping Identity product documentation is provided "as is" without warranty of any kind. Ping Identity disclaims all warranties, either express or implied, including the warranties of merchantability and fitness for a particular purpose. In no event shall Ping Identity or its suppliers be liable for any damages whatsoever including direct, indirect, incidental, consequential, loss of business profits or special damages, even if Ping Identity or its suppliers have been advised of the possibility of such damages. Some states do not allow the exclusion or limitation of liability for consequential or incidental damages so the foregoing limitation may not apply.

Table of Contents

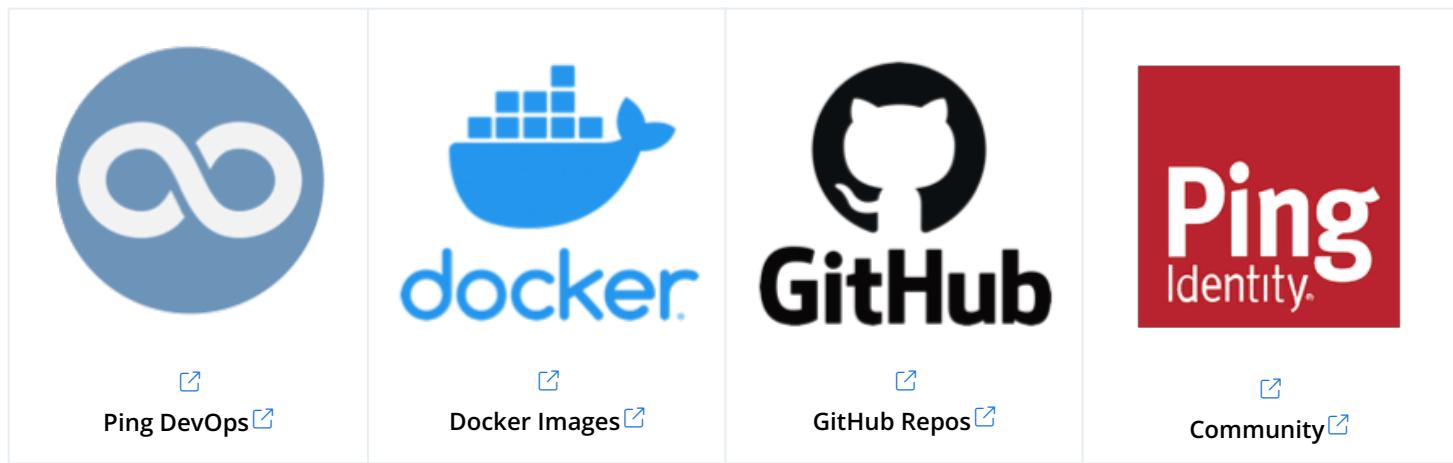
Get Started	
Welcome	4
Getting Started	4
Ideas or Suggestions.	7
Value Configs	
Introduction	9
Global	11
Container	11
External Image	12
Image	13
Ingress.	14
License	16
PrivateCert	17
Service.	18
Vault	19
VolumeMounts.	23
Workload	24
Supported Values	26
OpenShift Configuration	39
Examples	
Introduction	41
How To	
Update Tags	42
Release Notes	
Current Release	44
Previous Releases	44

Get Started

Welcome

This is the documentation for using [Helm](#) to deploy the Ping Identity container images. This single chart can be used to deploy any of the available [Ping Identity](#) products in a [Kubernetes](#) environment.

DevOps Resources



Prerequisites

- Kubernetes 1.16+
- Helm 3
- Ping Identity DevOps user/key

Adding the Helm Repo

```
helm repo add pingidentity https://docs.pingidentity.com/helm/
```

Removing the Repo

```
helm repo rm pingidentity
```

Getting Started

Helm [🔗](#) is a package deployment tool for Kubernetes [🔗](#). It can be used with Ping DevOps [🔗](#) to deploy all the components of the solution with a simple command.

Prerequisites

- Kubernetes Cluster
- Helm 3
- Ping Identity DevOps user/key

Note

Ping Helm charts support OpenShift. See [OpenShift Configuration](#) to learn how to configure the values.yaml file to do so.

Creating Ping DevOps Secret

The charts use a secret called `devops-secret` to obtain an evaluation license for running images.

- Eval License — Use your `PING_IDENTITY_DEVOPS_USER/PING_IDENTITY_DEVOPS_KEY` credentials along with your `PING_IDENTITY_ACCEPT_EULA` setting.
- For more information on obtaining credentials, click [here](#) [🔗](#).
- For more information on using the `pingctl` utility, click [here](#) [🔗](#).

```
pingctl k8s generate devops-secret | kubectl apply -f -
```

Installing Helm 3

Ensure that you have Helm 3 installed.

- Installing on macOS (or Linux with Brew):

```
brew install helm
```

- Installing on other OS:

<https://helm.sh/docs/intro/install/> [🔗](#)

Adding Helm Ping DevOps Repo

```
helm repo add pingidentity https://developer.pingidentity.com/helm/
```

Listing Ping DevOps Charts

```
helm search repo pingidentity
```

Updating Local Machine with the Latest Charts

```
helm repo update
```

Installing the Ping DevOps Chart

Install the `ping-devops` chart using the example below. In this case, it is installing a release called `pf`:

- PingFederate Admin instance
- PingFederate Engine instance

```
helm install pf pingidentity/ping-devops \
--set pingfederate-admin.enabled=true \
--set pingfederate-engine.enabled=true
```

or, if you have a `ping-devops-values.yaml` file:

```
# ping-devops-values.yaml
pingfederate-admin:
  enabled: true

pingfederate-engine:
  enabled: true
```

```
helm install pf pingidentity/ping-devops \
-f ping-devops-values.yaml
```

Accessing Deployments

By default, the components of a release are prefixed with the release name. Continuing this example, everything will be prefixed with `pf`. Use `kubectl` to see the pods created.

View Kubernetes resources installed:

```
# get just pods  
kubectl get pods --selector=app.kubernetes.io/instance=pf  
  
# or get even more  
kubectl get all --selector=app.kubernetes.io/instance=pf
```

View logs (from deployment):

```
kubectl logs deployment/pf-pingfederate-admin
```

Uninstalling Releases

To uninstall a release from Helm, use the following command (using `pf` as an example release):

```
helm uninstall pf
```

Contributing to the Ping Identity DevOps Program

Thanks for taking the time to help us improve our Helm chart!

You can contribute in various ways.

Reporting Bugs

Bugs are tracked as [GitHub issues](#). You can report a bug by submitting an issue in the project's issue tracker. To help the maintainers understand and reproduce the problem, please try to provide detailed information, including:

- A clear and descriptive title.
- A description of what happened and what you expected to happen.
- An example with the exact steps needed to reproduce the problem. If relevant, sample code is helpful.

Please understand that bug reports are reviewed and prioritized internally, and we may not be able to address all bug reports or provide an estimated time for resolution.

Suggesting Enhancements

As with bugs, requests are tracked as [GitHub issues](#). You can suggest an enhancement by submitting an issue in the project's issue tracker.

Please understand that enhancement requests are handled in the same way as bug reports, and we may not be able to address all enhancement requests or provide an estimated time for resolution.

i Note

If you would rather not have your issue discussed in public, you can email bug reports or enhancement requests to devops_program@pingidentity.com.

Contributing Code Changes

Ping Identity does not accept third-party code submissions.

Value Configs

Introduction

The charts make heavy use of `Values` yaml files to pass configuration details to the Helm Charts. As defined by [Helm Values Files](#), values are provided to the chart using the following mechanisms:

- `values.yaml` file in the chart
- Value files passed to Helm during install/upgrade with the `-f` flag
- Individual parameters passed with the `--set` flag

The list above is in order of specificity: `values.yaml` in the chart can be overridden with `-f` supplied files, which can in turn be overridden with the `--set` parameter.

The example below shows how values from the chart, a user-supplied `myconfig.yaml` file, and `--set` parameters are merged with each other to form merged values.

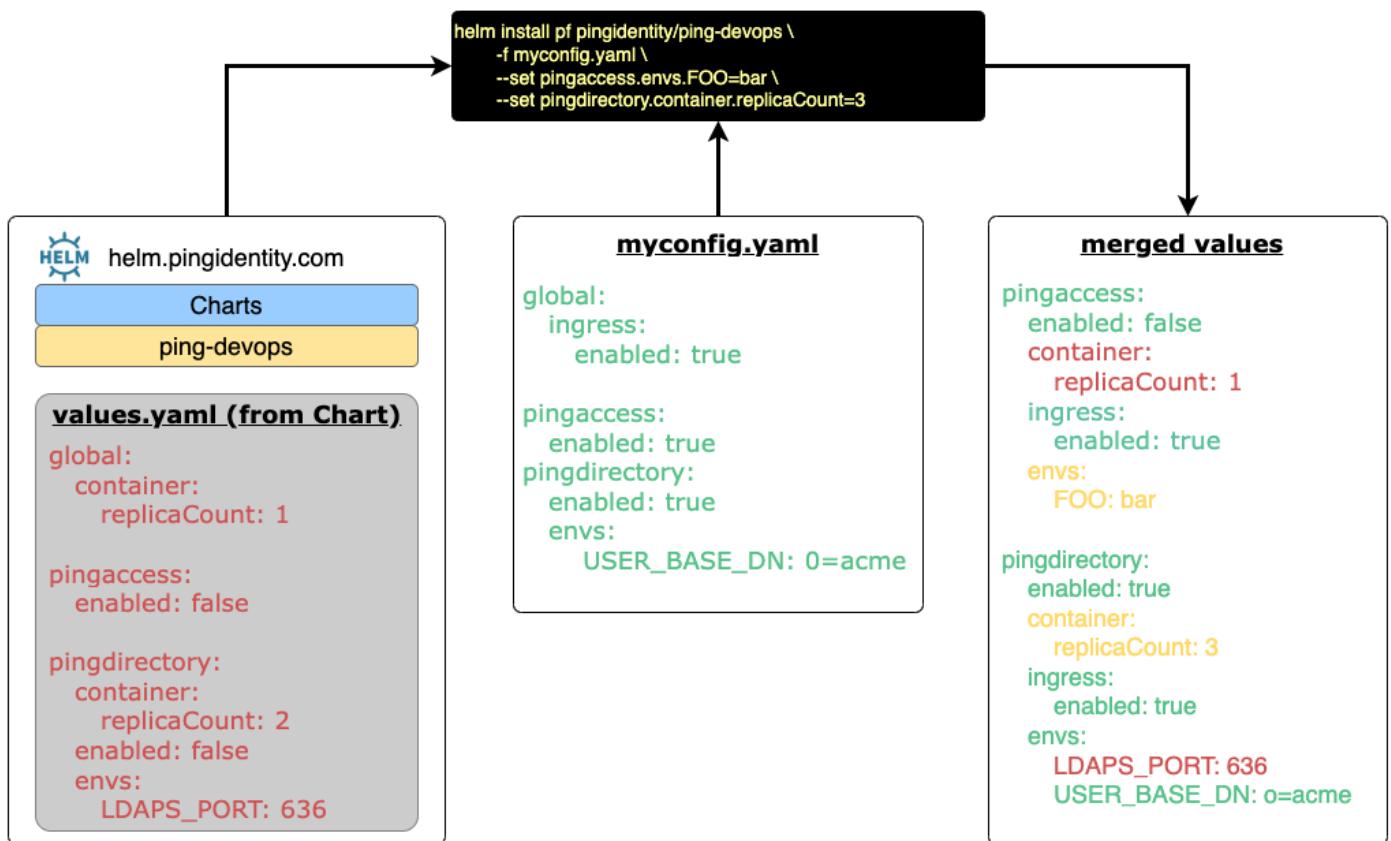


Chart Values

To see the values supplied by the chart, simply use the `helm show values` command to print them. This provides both the data as well as context-sensitive comments to each section.

```
helm show values pingidentity/ping-devops
#####
# Ping Identity DevOps values.yaml
#####
# ...
```

You can also see all the available values in the [helm-charts](#) repository on GitHub.

The default values are broken up into two major sections:

- **global** — Represents the base set of values that will be provided to each product section unless it's overridden in that section.
- **product** — For every image/product, the values will be merged with the global settings and take precedence.

Warning

Values can only be modified when merged. They cannot be deleted. Also, if a value is set to the boolean true, and merged with a boolean of false, it will always be true.

Global Section

The `global:` section of the values contains configuration that is available to each product section. If a value is set in `globals`, that will be available to every product. This is very powerful, as you can turn on the ingress for every product by simply setting:

```
global:
  ingress:
    enabled: true
```

This would in essence set `ingress.enabled=true` for every product:

```
pingaccess:
  ingress:
    enabled: true

pingdirectory:
  ingress:
    enable: true
```

and so on.

It is much easier to set something in the `global:` section rather than repeat it for each product. To enable the ingress for only a few specific products, leave the default of `global.ingress.enabled=false` and set that value for those product sections.

Product Sections

Just like the `global:` values, each product can have the same values as well as many more that are specific to that product/image. In the following example, persistent volume configuration is provided for PingDirectory:

```
pingdirectory:  
  persistentvolume:  
    enabled: true  
    volumes:  
      - name: out-dir  
        mountPath: /opt/out  
        storage: 8Gi  
        storageClassName:
```

global: Values

There is a top-level `global` value providing instructions on how to name all Kubernetes resources, so a deployer might deploy several releases under the same namespace.

addReleaseNameToResource

Provides global ability to add the Helm `.Release.Name` to Kubernetes resources.

Value	Description	Example: (Release.Name=acme, resource=pingdirectory)
prepend	Prepends Release.Name [DEFAULT]	acme-pingdirectory
append	Appends Release.Name	pingdirectory-acme
none	No use of Release.Name	pingdirectory

Container Configuration

[Kubernetes Workload Controller](#) resources are created depending on configuration values:

- [Deployments](#)
- [StatefulSets](#)

Global Section

Default yaml defined in the global `container` section:

```
global:  
  container:  
    replicaCount: 1  
    resources:  
      requests:  
        cpu: 0  
        memory: 0  
      limits:  
        cpu: 0  
        memory: 0  
    nodeSelector: {}  
    tolerations: []  
    affinity: {}  
    terminationGracePeriodSeconds: 30  
    envFrom: []  
    lifecycle: {}  
    probes:  
      livenessProbe:  
        exec:  
          command:  
            - /opt/liveness.sh  
        initialDelaySeconds: 30  
        periodSeconds: 30  
        timeoutSeconds: 5  
        successThreshold: 1  
        failureThreshold: 4  
      readinessProbe:  
        exec:  
          command:  
            - /opt/readiness.sh  
        initialDelaySeconds: 30  
        periodSeconds: 5  
        timeoutSeconds: 5  
        successThreshold: 1  
        failureThreshold: 4  
      startupProbe:  
        exec:  
          command:  
            - /opt/liveness.sh  
        periodSeconds: 10  
        timeoutSeconds: 5  
        failureThreshold: 90
```

Probes Configuration

Kubernetes Probes [🔗](#) defined in the `container:` section will be added to workloads (that is, Deployments/StatefulSets).

Fields used to configure probes can be found in the [Kubernetes documentation](#) [🔗](#).

External Image Configuration

Defines an external image for initContainer utilities.

Global Section

Default yaml defined in the global `externalImage` section:

```
global:  
  externalImage:  
    pingtoolkit: pingidentity/pingtoolkit:latest
```

External Image Parameters	Description
pingtoolkit	Registry, image and tag location for pingtoolkit. Used for primarily during init containers.

Note

If your Kubernetes cluster doesn't have access to an external Docker repository, you can download and save the pingtoolkit image to your local repo. Setting this to your local repo will cause the charts to use that image.

Image Configuration

Provides values to define kubernetes image information to deployments and statefulsets.

Global Section

Default image yaml defined in the global section:

```
global:  
  image:  
    repository: pingidentity  
    name:          # Set in product section  
    tag: 2307  
    pullPolicy: Always  
  imagePullSecrets: []  # As needed for authentication to private repositories  
  # - name: myregkeysecretname
```

Product Section

Each product section specifies the name by default.

```
pingaccess-admin:  
  image:  
    name: pingaccess
```

To have images use a different repository and tag, use the following:

```
global:
  image:
    tag: edge
  repository: my.company.docker-repo.com
```

This snippet would result in pulling a PingAccess image from `my.company.docker-repo.com/pingaccess:edge`.

Ingress Configuration

[Kubernetes Ingress resources](#) are created depending on configuration values.

Global Section

Default yaml defined in the global `ingress` section, followed by definitions for each parameter:

```
global:
  ingress:
    enabled: false
    addReleaseNameToHost: subdomain
    defaultDomain: example.com
    defaultTlsSecret:
    annotations: {}
    spec: {}
```

Ingress Parameters	Description	Options	Default Value
enabled	Enables ingress definition.		false
addReleaseNameToHost	How <code>helm release-name</code> should be added to host.	prepend append subdomain none	subdomain
defaultDomain	Default DNS domain to use. Replaces the string <code>_defaultDomain_</code> .		
defaultTlsSecret	Default TLS Secret to use. Replaces the string <code>_defaultTlsSecret_</code> .		example.com
annotations	Annotations are used to provide configuration details to specific ingress controller types.	* see option for nginx ingress	{}
spec.ingressClassName	This value is replacing the <code>kubernetes.io/ingress.class</code> annotation. See this page for details.	name of the IngressClass resource	{}

Annotations example for nginx ingress:

```
annotations:
  nginx.ingress.kubernetes.io/backend-protocol: "HTTPS"
```

ingressClassName specification example for nginx ingress:

```
spec:
  # Must match the name of the IngressClass resource
  ingressClassName: nginx-public
```

Product Section

Default yaml defined in the product `ingress` section, followed by definitions for each parameter:

```
ingress:
  hosts:
    - host: pingfederate-admin._defaultDomain_
      paths:
        - path: /
          backend:
            serviceName: admin
  tls:
    - secretName: _defaultTlsSecret_
      hosts:
```

Ingress Parameters	Description	Default Value
hosts	Array of hosts definitions	
hosts[].host	Full DNS name of host to use for external name. "_defaultDomain_" will be replaced with .defaultDomain	{product- name}._defaultDomain_
hosts[].paths	Array of paths to define for host	
hosts[].paths[].path	Path on external ingress	
hosts[].paths[].backend.serviceName	Name of the service to map to. This will result in the ingressPort on the server to be used.	
tls	Array of tls definitions	
tls[].secretName	Certificate secret to use	_defaultTlsSecret_
tls[].hosts	Array of specific hosts	

Example use of `_defaultDomain_` and `addReleaseNameToHost`:

```
helm ReleaseName = acme
  defaultDomain = example.com
  addReleaseNameToHost = subdomain
ingress.hosts[0].host = pingfed-admin._defaultDomain_

Resulting host will be: pingfed-admin.acme.example.com
                           ^          ^^^^^^
                           |          |
  ReleaseName           defaultDomain
```

Example Ingress Manifest

Below is an example product ingress for `pingfederate-admin` when deployed by Helm with a release-name of `acme`. It includes an ingress for the admin service (9999) using the default domain and tls secret, defined in the global section (if set).

```
kind: Ingress
metadata:
  annotations:
    ...
spec:
  rules:
    - host: pingfederate-admin.acme.example.com
      http:
        paths:
          - backend:
              serviceName: acme-pingfederate-admin
              serviceName: admin
              path: /
      tls:
        - hosts:
          - pingfederate-admin.acme.example.com
        secretName: ""
```

License Configuration

Provides a secret used for obtaining evaluation licenses for Ping Identity products.

Global Section

Default yaml defined in the global `license` section, followed by definitions for each parameter:

```
global:
  license:
  secret:
    dev0ps: devops-secret
```

License Parameters	Description	Default Value
secret.devops	Secret containing PING_IDENTITY_DEVOPS_USER/KEY values.	devops-secret

Note

Use the `pingctl` command-line tool to create the devops-secret with your Ping Identity DevOps user and key.
`pingctl kubernetes generate devops-secret | kubectl apply -f`

PrivateCert Configuration

Generates a private certificate (.crt and .key) based on the internal hostname of the service.

Global Section

Note

`privateCert` is currently only supported by PingAccess.

Default yaml defined in the global `privateCert` section. By default, certificates will not be generated. It is advised to NOT generate internal certs at the global level, as many services don't need a private cert on the internal service.

```
global:
#####
# Internal Certificates
#
# If set to true, then an internal certificate secret will
# be created along with mount of the certificate in
# /run/secrets/internal-cert (creates a tls.crt and tls.key)
#
# By default the Issuer of the cert will be the service name
# created by the Helm Chart. Additionally, the ingress hosts,
# if enabled, will be added to the list of X509v3 Subject Alternative Name
#
# Use the additionalHosts and additionalIPs if additional custom
# names and ips are needed.
#
#     privateCert.generate: {true | false}
#     privateCert.additionalHosts: {optional array of hosts}
#     privateCert.additionalIPs: {optional array of IP Addresses}
#####
privateCert:
  generate: false
  additionalHosts: []
  additionalIPs: []
```

Product Section

Generating an internal certificate is as simple as setting `privateCert.generate` to true.

Here's an example of generating an internal certificate for `pingaccess-engine`:

```
pingaccess-admin:  
  privateCert:  
    generate:true
```

This will ultimately create a secret named `{release-productname}-private-cert` containing a valid `tls.crt` and `tls.key`.

By default, the issuer of the cert will be the service name created by the Helm Chart. Additionally, the ingress hosts, if enabled, will be added to the list of `X509v3 Subject Alternative Name`.

The product image will then create an init container to generate a pkcs12 file that will be placed in `/run/secrets/private-keystore/keystore.env`, which will be mounted into the running container.

When the container's hooks are running, it will source the environment variables in this `keystore.env`. The default variables set are:

- `PRIVATE_KEYSTORE_PIN={base64 random pin}`
- `PRIVATE_KEYSTORE_TYPE=pkcs12`
- `PRIVATE_KEYSTORE={pkcs12 keystore}`

These environment variables are required in the `data.jsonsubst` file in order to use the generated `privateCert`. They can be used in any server-profile artifacts to be replaced when the images are started.

Service Configuration

[Kubernetes Service resources](#) are created depending on configuration values.

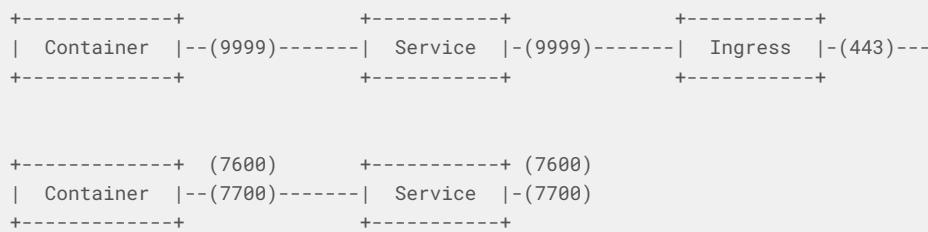
Product Section

Default yaml defined in the product `services` section. The example found in the `pingfederate-admin` section is:

```
services:  
  admin:  
    servicePort: 9999  
    containerPort: 9999  
    ingressPort: 443  
    dataService: true  
  clusterbind:  
    servicePort: 7600  
    containerPort: 7600  
    clusterService: true  
  clusterfail:  
    servicePort: 7700  
    containerPort: 7700  
    clusterService: true  
  clusterExternalDNSHostname:
```

Service Parameters	Description
services	Array of services
services[].{name}	Service Name. (i.e. https, ldap, admin, api)
services[].{name}.servicePort	External port of service
services[].{name}.containerPort	Port on target container
services[].{name}.ingressPort	Port on ingress container (if ingress is used)
services[].{name}.dataService	Adds to a ClusterIP service with single DNS/IP
services[].{name}.clusterService	Adds to a headless service with DNS request returning all IPs
services.clusterExternalDNSHostname	

The example above will create a container/service/ingress that looks like this:



Vault Configuration

The current Helm chart support is provided for Hashicorp Vault annotations and use of the Hashicorp injector. More information on Hashicorp Vault annotations can be found [here](#).

Note

The PingIdentity DevOps images and Helm chart only support version 2 of the KV secrets engine API for Vault secrets. PingDirectory itself currently only supports KV version 1 for password storage schemes. Learn more in the [Vault KV secrets engine documentation](#).

Vault Secret Values

An example vault values section looks like:

```

vault:
  enabled: true
  hashicorp:
    annotations:
      role: {hashicorp-vault-role}
    secretPrefix: {path to secret}
    secrets:
      {secret-name}:
        {secret-key | to-json}:
          path: /opt/in/some/location/secrets
          file: devops-secret.env

```

The `vault.hashicorp.secrets` is a map that specifies each `secret` to pull from the vault. And for each secret, a map specifies the `key` to pull with instructions of where to place the secret based on `path` and `file`.

License Parameters	Description	Default Value
secrets.{secret}	map of secret	devops-secret
secrets.{secret}.{key}	map of key	pingaccess.lic
secrets.{secret}.{key}.path	optional: location of secret. Defaults to vault.annotation.secret-volume-path	/opt/in/some/path
secrets.{secret}.{key}.file	required: file name secrets placed into	pingaccess.lic

Special Key Name (to-json)

There is a special key name that can be provided that will drop the raw secret into the container as its json representation with all the secret key names/values.

If dropped into the `SECRETS_DIR` (defaults to `/run/secrets`) directory, these files will be processed as:

- `PROPERTY_FILE` if the file ends in `.env`, or
- Separate files will be created for each key=value pair.

See the example below in this document for the transformation that occurs with the `devops-secret.env`.

Vault Annotations

Default yaml defined in the global `vault` section. The options of annotation names/values can be found at [vault definitions](#).

For each of the annotations, the helm chart will automatically pre-pend the annotation with the hashicorp annotation prefix of `vault.hashicorp.com`. See example below.

```
global:  
  vault:  
    enabled: false  
    hashicorp:  
      annotations:  
        agent-inject: true  
        agent-init-first: true  
        agent-pre-populate-only: true  
        log-level: info  
        preserve-secret-case: true  
        role: k8s-default  
        secret-volume-path: /run/secrets
```

The serviceAccount used by Vault will match the default serviceAccount for the workload.

Example

The following includes an example Hashicorp Vault secrets as well as a value values.yaml that make use of the secrets and an example of where secrets will be placed into the container.

Example: Hashicorp Vault secrets

```
SECRET:secrets/jsmith@example.com/jsmith-namespace/licenses  
  
{  
  "pingaccess-6.2": "Product=PingAccess\\nVersion=6.2...","  
  "pingdirectory-8.2": "Product=PingDirectory\\nVersion=8.2...","  
  "pingfederate-10.2": "Product=PingFederate\\nVersion=10.2..."  
}  
  
SECRET: secrets/jsmith@example.com/jsmith-namespace/devops-secrets.env  
  
{  
  "PING_IDENTITY_ACCEPT_EULA": "YES",  
  "PING_IDENTITY_DEVOPS_KEY": "d254.....-....-....-.....",  
  "PING_IDENTITY_DEVOPS_USER": "jsmith@example.com"  
}  
  
SECRET: secrets/jsmith@example.com/jsmith-namespace/certs  
  
{  
  "tls.crt": "LS0tLS1CRUdJ...a9dk",  
  "tls.key": "LS0tLS1CRUdJ...38sj"  
}
```

Example: Vault secrets .yaml

```
pingfederate-admin:  
  vault:  
    hashicorp:  
      secrets:  
        devops-secret.env:  
          to-json:  
            file: devops-secret.env  
        licenses:  
          pingaccess-6.2:  
            file: pingaccess.lic  
            path: /opt/in/some/location/licenses  
        test-certs:  
          to-json:  
            file: test-certs
```

Places the following files into the container.

Example: Container files

FILE: /run/secrets/devops-secret.env

```
PING_IDENTITY_ACCEPT_EULA="YES"  
PING_IDENTITY_DEVOPS_KEY="d254....-....-....-....-....-...."  
PING_IDENTITY_DEVOPS_USER="jsmith@example.com"
```

FILE: /opt/in/some/location/licenses/pingaccess.lic

```
Product=PingAccess  
Version=6.2  
...
```

FILE: /run/secrets/tls.crt

```
LS0tLS1CRUdj...a9dk
```

FILE: /run/secrets/tls.key

```
LS0tLS1CRUdj...38sj
```

Using Vault Secrets to Mount base64-Encoded Keystore Files

To pull keystore files from a Vault cluster, create a secret with separate keys for each individual keystore file. The value of each key should be the base64 representation of the file that needs to be mounted. The names of the keys as well as the name of the secret will be used in the Helm values yaml when mounting the keystore files.

Environment variables or other configuration for the product being deployed will need to be set to point to the location where the keystores are being mounted. For example, for PingDirectory:

```
KEYSTORE_FILE=/run/secrets/mykeystore.jks  
KEYSTORE_PIN_FILE=/run/secrets/mykeystore.pin
```

Ensure the Vault cluster is accessible to the Kubernetes cluster where your Helm release is being deployed. You can then use the Vault annotations to mount the keystore files in the desired location. For an example, see the "Vault Keystores" example on the [Helm examples page](#).

VolumeMounts Configuration

Provides support for mounting secret or configMap volumes on a workload container.

Global/Product Section

Adds ability to use **secret** and **configMap** data in a container via a VolumeMount. A common use for this configuration is bringing product licenses or scripts into the container.

Example of creating two volume mounts in container from secret and configMap:

```
pingfederate-admin:  
  enabled: true  
  volumes:  
    - name: pf-props  
      configMap:  
        name: pingfederate-props  
    - name: pf-license  
      secret:  
        secretName: pingfederate-license  
  volumeMounts:  
    - mountPath: /opt/in/etc/pingfederate.properties  
      name: pf-props  
    - mountPath: /opt/in/instance/server/default/conf/pingfederate.lic  
      name: pf-license
```



Note

[Secrets](#) and [ConfigMaps](#) must be created in the cluster prior to deploying the helm chart.

In this case, a secret (called `pingfederate-license`) and configMap (called `pingfederate-props`) will bring in a couple of key values (license, hello) and (pf-props) into the container as specific files. The resulting object will look like this:

Example of kubectl describe of pingfederate-admin container

```
Containers:  
pingfederate-admin:  
  Mounts:  
    /opt/in/etc/pingfederate.properties from pingfederate-props (ro,path="pingfederate.properties")  
    /opt/in/instance/server/default/conf/pingfederate.lic from pingfederate-license (ro,path="pingfederate.lic")  
Volumes:  
pingfederate-license:  
  Type: Secret (a volume populated by a Secret)  
  SecretName: pingfederate-license  
  Optional: false  
pingfederate-props:  
  Type: ConfigMap (a volume populated by a ConfigMap)  
  Name: pingfederate-props  
  Optional: false
```

Workload Configuration

Kubernetes Workload resources are created depending on configuration values:

- [Deployments](#)
- [StatefulSets](#)

Global Section

Default yaml is defined in the global `workload` section. Individual products override these defaults based on the required workload.

```

global:
  workload:
    type: Deployment

  deployment:
    strategy:
      type: RollingUpdate
      rollingUpdate:
        maxSurge: 1
        maxUnavailable: 0

  statefulSet:
    partition: 0

  persistentvolume:
    enabled: true
    volumes:
      out-dir:
        mountPath: /opt/out
        persistentVolumeClaim:
          accessModes:
            - ReadWriteOnce
          storageClassName:
          resources:
            requests:
              storage: 4Gi

  securityContext:
    fsGroup: 9999
  securityContext: {}

```

Workload Parameters	Description
type	One of Deployment or StatefulSet
deployment.strategy.type	One of RollingUpdate or ReCreate
deployment.strategy.rollingUpdate	If type=RollingUpdate
statefulSet.partition	Used for canary testing if n>0
statefulSet.persistentVolume	Provides details around creation of PVC/Volumes (see below)
securityContext	Provides security context details for starting container as different user/group (see below)
securityContext.fsGroup	Sets the group id on fileSystem writes. This is needed especially for mounted volumes (pvs)

Persistent Volumes

For every volume defined in the volumes list, three items will be created in the StatefulSet:

- container.volumeMounts — name and mountPath
- template.spec.volume — name and persistentVolumeClaim.claimName
- spec.volumeClaimTemplates — persistentVolumeClaim

For further details, see the [Kubernetes documentation](#).

Security Context

To run the containers with a different user/group/fsGroup, use the following example to set those details on the deployment/statefulset:

```
global:  
  workload:  
    container:  
      securityContext:  
        runAsGroup: 9999  
        runAsUser: 9031  
        fsGroup: 9999
```

WaitFor

For each product, you can provide a `waitFor` structure indicating the name, service, and timeout (in seconds) for which the container should wait before continuing (default: 300). This setting will inject an initContainer using the PingToolkit wait-for utility that relies on `nc host:port` before continuing.

Example: PingFederate Admin waiting on pingdirectory ldaps service to be available

```
pingfederate-admin:  
  container:  
    waitFor:  
      pingdirectory:  
        service: ldaps  
        timeoutSeconds: 600  
      pingauthorize:  
        service: https  
        timeoutSeconds: 300
```

- By default, the pingfederate-engine will waitFor pingfederate-admin before it starts.
- By default, the pingaccess-engine will waitFor pingaccess-admin before it starts.

List of Supported Values

These are the values supported in the ping-devops chart. In general, values specified in the global section can be overridden for individual products. The product sections have many global fields overridden by default (workloads, services, etc.).

Global Values

Name	Description	Default
global.annotations	Annotations listed, will be added to all Kubernetes resources.	{}
global.labels	Labels listed, will be added to all Kubernetes resources.	{}
global.envs	Environment variables listed will be added to the global-env-vars configmap	{}
global.addReleaseNameToResource	Provides global ability to add names to kubernetes resources. One of {none, append, prepend}	prepend
global.ingress.enabled		false
global.ingress.addReleaseToHost	Add release to host. One of {prepend, append, subdomain, none}	subdomain
global.ingress.defaultDomain	Replaces with " <i>defaultDomain</i> " in host fields	example.com
global.ingress.defaultTlsSecret	Replaces with " <i>defaultTlsSecret</i> " in tls.secretName	
global.ingress.annotations		{}
global.ingress.spec.ingressClassName		
global.privateCert.generate	If true, then an internal certificate secret will be created along with mount of the certificate in /run/secrets/internal-cert (creates a tls.crt and tls.key). By default the Issuer of the cert will be the service name created by the Helm Chart. Additionally, the ingress hosts, if enabled, will be added to the list of X509v3 Subject Alternative Name	false
global.privateCert.format	The format of the certificate to be generated. Used "pingaccess-fips-pem" to generate a valid certificate for running PingAccess in FIPS mode. Any other value will generate a PKCS12 keystore with the generated certificate.	PKCS12
global.privateCert.additionalHosts	Additional hosts for the cert	[]
global.privateCert.additionalIPs	Additional IP addresses for the cert	[]
global.masterPassword	Uses Helm function derivePassword, which uses the master password specification: https://masterpassword.app/masterpassword-algorithm.pdf	
global.masterPassword.enabled	Enable master password	false

Name	Description	Default
global.masterPassword.strength	Master password template. One of {long, maximum}	
global.masterPassword.name	Defaults to release name	
global.masterPassword.site	Defaults to chart name	
global.masterPassword.secret	Defaults to release namespace	
global.vault	Hashicorp Vault configuration	
global.vault.enabled	Enable Vault	false
global.vault.hashicorp.annotations	Annotation names, which will be appended to 'vault.hashicorp.com/' in the annotation. The vault.hashicorp.annotations.serviceAccountName value will be overwritten by the service account generated for the workload if there is one.	
global.vault.secretPrefix	Prefix that will be prepended to any secrets being injected.	""
global.vault.secrets	Vault secrets to pull in	{}
global.imagePullSecrets	Repository authentication using secret defined as a docker-registry secret in Kubernetes.	[]
global.image.repository	Default image registry is not the fully-qualified name of the image Example: image.repository: pingidentity, docker.io, 123.dkr.ecr.us-west-1.amazonaws.com	pingidentity
global.image.repositoryFqn	Docker image repository fully-qualified name. Overrides image.repository and image.name on the pod image spec Example: image.repositoryFqn: pingidentity/pingfederate, docker.io/my-pingfederate	
global.image.name	Default image name MUST be set in child chart Example: image.name: pingfederate	
global.image.tag	Default image tag	2505
global.image.pullPolicy	Default image pull policy	IfNotPresent
global.rbac.generateServiceAccount	Set to true to generate a service account for the workload.	false

Name	Description	Default
global.rbac.serviceAccountName	Name of the service account that will be generated. The default value of "defaultServiceAccountName" will result in a service account named based on the Helm installation and the specific workload being deployed. If generateServiceAccount and generateGlobalServiceAccount are false, this value can also refer to a service account created outside of Helm.	defaultServiceAccountName
global.rbac.generateRoleAndRoleBinding	Set to true to generate a Role and RoleBinding corresponding to the workload service account.	false
global.rbac.generateGlobalServiceAccount	Set to true to generate a service account for the entire installation. This global service account will be used for workloads that do not generate their own service account.	false
global.rbac.generateGlobalRoleAndRoleBinding	Set to true to generate a Role and RoleBinding corresponding to the global service account for the entire installation.	false
global.rbac.applyServiceAccountToWorkload	Set to true (the default) to apply to service account to the workload.	true
global.rbac.role	This yaml will be directly inserted into the generated Role when generateRoleAndRoleBinding and/or generateGlobalRoleAndRoleBinding are true. The rules for the Role can be set here.	get, watch, and list verbs for the pods resource
global.rbac.serviceAccountAnnotations	Any custom annotations to add to the service account.	
global.rbac.roleAnnotations	Any custom annotations to add to the role.	
global.rbac.roleBindingAnnotations	Any custom annotations to add to the role binding.	
global.rbac.serviceAccountLabels	Any custom labels to add to the service account.	
global.rbac.roleLabels	Any custom labels to add to the role.	
global.rbac.roleBindingLabels	Any custom labels to add to the role binding.	

Name	Description	Default
<code>global.externalImage</code>	Provides ability to use external images for various purposes such as using curl, waitfor, etc. A pingtoolkit image is included by default for running waitFor and generating private cert initContainers. A pingaccess image is also included by default to allow generating an encrypted PEM-formatted cert that is compatible with FIPS mode. Any values specified on the image will be copied directly to the k8s spec for the container, except for the <code>externallImage.{name}.image</code> section, which follows the format of the <code>global.image</code> section. If no image section is specified (the default), the corresponding value from the product values section will be used. For example, if <code>externallImage.pingtoolkit.image</code> is empty, the values from the top-level <code>pingtoolkit.image</code> section will be used.	<code>{pingtoolkit, pingaccess}</code>
<code>global.services</code>	Services mapping a port to a targetPort on the corresponding container	<code>{}</code>
<code>global.services.clusterExternalDNSHostname</code>	Value for the <code>external-dns.alpha.kubernetes.io/hostname</code> annotation for the cluster service.	
<code>global.services.clusterServiceName</code>	If set, then this name will be used as the cluster service name (i.e <code>clusterService == true</code>).	
<code>global.services.useLoadBalancerForDataService</code>	If true, the data service will be created with type: LoadBalancer.	<code>false</code>
<code>global.services.serviceName.dataService</code>	If true, a ClusterIP service is created reachable within the cluster. A single IP is provided and the service will round-robin across the backend containers	
<code>global.services.serviceName.clusterService</code>	If true, a headless service is created, explicitly specifying "None" for the clusterIP. DNS requests to this service will provide one of the IPs of the backend containers	
<code>global.services.serviceName.containerPort</code>	Port on the kubernetes container	
<code>global.services.serviceName.servicePort</code>	Port available from the kubernetes service. If <code>clusterService=true</code> this port on the cluster service is not really used, as the headless service always maps through to the container port	
<code>global.services.serviceName.ingressPort</code>	Port available from the kubernetes ingress	
<code>global.services.annotations</code>	Any custom annotations to add to the service.	

Name	Description	Default
global.services.clusterServiceAnnotations	Any custom annotations to add to the ClusterIP service.	
global.services.labels	Any custom labels to add to the service.	
global.services.clusterServiceLabels	Any custom labels to add to the ClusterIP service.	

Workload Values – Deployment and StatefulSet

Name	Description	Default
global.workload	Can be Deployment or StatefulSet	Deployment
global.workload.annotations	Annotations to apply to the template in the workload. To apply top-level annotations to the Deployment or StatefulSet itself, use global.workload.deployment.annotations or global.workload.statefulSet.annotations.	
global.workload.labels	Labels to apply to the template in the workload. To apply top-level labels to the Deployment or StatefulSet itself, use global.workload.deployment.labels or global.workload.statefulSet.labels.	
global.workload.schedulerName	K8s scheduler	default-scheduler
global.workload.shareProcessNamespace	Set shareProcessNamespace in the pod spec	false
global.workload.enableServiceLinks	indicates whether info about services can be added as env variables	true
global.workload.topologySpreadConstraints	Configuration of pod spread across cluster zones	[]
global.workload.deployment	Deployment workload configuration	
global.workload.deployment.strategy	Deployment pod replacement strategy	
global.workload.deployment.strategy.type	Strategy type	RollingUpdate
global.workload.deployment.strategy.rollingUpdate.maxSurge	Max surge, only applicable for RollingUpdate type	1

Name	Description	Default
global.workload.deployment.strategy.rollingUpdate.maxUnavailable	Max unavailable, only applicable for RollingUpdate type	0
global.workload.deployment.annotations	Annotations to apply to the top-level Deployment. To apply annotations to the template within the Deployment, use global.workload.annotations.	
global.workload.deployment.labels	Labels to apply to the top-level Deployment. To apply labels to the template within the Deployment, use global.workload.labels.	
global.workload.statefulSet	StatefulSet workload configuration	
global.workload.statefulSet.partition	Used for canary testing if n>0	0
global.workload.statefulSet.persistentvolume.enabled	Enable persistent volumes	true
global.workload.statefulSet.persistentvolume.volumes	For every volume defined in the volumes list, 3 items will be created in the StatefulSet: 1. container.volumeMounts - name and mountPath. 2. template.spec.volume - name and persistentVolumeClaim.claimName. 3. spec.volumeClaimTemplates - persistentVolumeClaim.	{out-dir}
global.workload.statefulSet.persistentvolume.volumes.volumeName.mountPath	Mount path for the volume	
global.workload.statefulSet.persistentvolume.volumes.volumeName.persistentVolumeClaim	volumeClaimTemplate	
global.workload.statefulSet.podManagementPolicy	Controls how pods are created during initial scale up, when replacing pods on nodes, or when scaling down. The default behavior is OrderedReady. The Parallel podManagementPolicy allows for starting up and scaling down multiple Pods simultaneously. Updates are not affected. The only products that support Parallel are PingDirectory and PingDataSync, on versions 2209 and later. When using the Parallel policy, consider setting the RETRY_TIMEOUT_SECONDS environment variable to a higher value (it defaults to 180) for the Pods. If the value is too low with many servers starting at once, it may lead to some Pods restarting unnecessarily during the initial workload startup.	OrderedReady

Name	Description	Default
global.workload.statefulSet.annotations	Annotations to apply to the top-level StatefulSet. To apply annotations to the template within the StatefulSet, use global.workload.annotations.	
global.workload.statefulSet.labels	Labels to apply to the top-level StatefulSet. To apply labels to the template within the StatefulSet, use global.workload.labels.	
global.workload.securityContext	securityContext for the workload Pod spec. The securityContext defined will be inserted directly into the Pod spec. The user (9031) and group (0) represent the current user and group used with PingIdentity images (except PingDelegator). The fsGroup is required for any workloads that volumeMount a pvc (i.e. StatefulSets). Set as securityContext: null when no generated securityContext is desired.	fsGroup 0, runAsUser 9031, runAsGroup 0
global.clustering.autoscaling	Configure Horizontal Pod Autoscaling	
global.clustering.autoscaling.enabled	Enable Horizontal Pod Autoscaling. If enabled, ensure that proper container.resources values are set and coordinated with the targetCPUUtilizationPercentage or targetMemoryUtilizationPercentage	false
global.clustering.autoscaling.minReplicas	Autoscaler minimum replicas	1
global.clustering.autoscaling.maxReplicas	Autoscaler maximum replicas	4
global.clustering.autoscaling.targetCPUUtilizationPercentage	Target CPU utilization	75
global.clustering.autoscaling.targetMemoryUtilizationPercentage	Target memory utilization	
global.clustering.autoscaling.annotations	Custom annotations for the HPA.	
global.clustering.autoscaling.labels	Custom labels for the HPA.	
global.clustering.autoscaling.behavior	Custom HPA behavior yaml	{}
global.clustering.autoscalingMetricsTemplate	Custom HPA metrics yaml	[]

Name	Description	Default
global.container	Configure the container in the workload Pod spec	
global.workload.container.securityContext	securityContext at the container level for the workload. The securityContext defined will be inserted directly into the spec for the main container of the Pod. Container-level securityContext values will overwrite any corresponding values from the Pod-level securityContext.	allowPrivilegeEscalation: false, capabilities: drop: ALL
global.container.replicaCount	Number of replicas for workload	1
global.container.resources	container resources yaml to insert into Pod spec	
global.container.nodeSelector	nodeSelector yaml to insert into Pod spec	{}
global.container.tolerations	tolerations yaml to insert into Pod spec	[]
global.container.affinity	affinity yaml to insert into Pod spec	{}
global.container.terminationGracePeriodSeconds	termination grace period	30
global.container.envFrom	envFrom yaml to insert into Pod spec	[]
global.container.env	Additional environment variables to insert into the Pod spec. Unlike the global.envs values, these will be set directly on the Pod. global.envs values are set in ConfigMaps rather than on the Pod directly. This value allows for setting the valueFrom field for an environment variable when necessary.	[]
global.container.lifecycle	lifecycle yaml to insert into Pod spec	
global.container.probes	probes yaml to insert into Pod spec	liveness, readiness, and startup probes defined

Other Global Defaults

Name	Description	Default
global.license.secret.devops	Identify the k8s secret containing the DevOps USER/KEY if used during deployment. pingctl can be used to generate the devops-secret	devops-secret

Name	Description	Default
global.utilitySidecar	Deploy a utility sidecar for running command-line tools. This sidecar is useful for command line utilities like collect-support-data. The sidecar will remain running alongside the workload, even when the sidecar isn't being used. It does not need to be listed in the includeSidecars value.	
global.utilitySidecar.enabled	Enable the utility sidecar	false
global.utilitySidecar.resources	Set k8s resources yaml for the sidecar spec	1 CPU and 2g memory limit, 0 CPU and 128Mi memory request
global.utilitySidecar.env	Environment variables for the sidecar	
global.utilitySidecar.securityContext	securityContext at the container level for the sidecar. The securityContext defined will be inserted directly into the spec for the sidecar. By default no container securityContext is defined. In Kubernetes when a container-level securityContext is set, it will overwrite any corresponding values from the Pod-level securityContext.	allowPrivilegeEscalation: false, capabilities: drop: ALL
global.includeSidecars	names of sidecars to include, from the top-level <code>sidecars</code> value	[]
global.includeInitContainers	names of sidecars to include, from the top-level <code>initContainers</code> value	[]
global.includeVolumes	names of sidecars to include, from the top-level <code>volumes</code> value	[]

Shared Utilities

Name	Description	Default
sidecars	Sidecar yaml definitions available to product workload spec	{}
initContainers	initContainer yaml definitions available to product workload spec	{}
volumes	volume yaml definitions available to product workload spec for sidecars, initContainers, or main product containers	{}
configMaps	configMap yaml definitions available to product workload spec for sidecars or main product containers	{}

Image/Product Values

Name	Description	Default
ldap-sdk-tools	LDAP SDK tools values	
ldap-sdk-tools.enabled	Enable LDAP SDK tools deployment	false
pingfederate-admin	PingFederate admin values	
pingfederate-admin.enabled	Enable PingFederate admin deployment	false
pingfederate-admin.cronjob	CronJobs run a kubectl exec command to run commands on a utility sidecar container. They will also create the necessary ServiceAccount, Role, and RoleBinding to run the jobs	
pingfederate-admin.cronjob.enabled	Enable the PingFederate Admin CronJob	false
pingfederate-admin.cronjob.spec	yaml to insert into the created CronJob spec. If specified, this will override any other specified values for the CronJob.	
pingfederate-admin.cronjob.spec.jobTemplate	yaml to override default jobTemplate. If a jobTemplate is not overridden, a default template will be inserted.	
pingfederate-admin.cronjob.image	Image to run the Jobs. The image must include kubectl	bitname/ kubectl:latest
pingfederate-admin.cronjob.args	Job arguments	[]
pingfederate-admin.cronjob.podSecurityContext	securityContext for the pod in the jobTemplate. This will be used if a jobTemplate is not specified.	null
pingfederate-admin.cronjob.containerSecurityContext	securityContext for the container in the jobTemplate. This will be used if a jobTemplate is not specified.	allowPrivilegeEscalation: false, capabilities: drop: ALL
pingfederate-engine	PingFederate engine values	
pingfederate-engine.enabled	Enable PingFederate engine deployment	false
pingfederate-engine.cronjob	CronJobs run a kubectl exec command to run commands on a utility sidecar container. They will also create the necessary ServiceAccount, Role, and RoleBinding to run the jobs	
pingfederate-engine.cronjob.enabled	Enable the PingFederate engine CronJob	false

Name	Description	Default
<code>pingfederate-engine.cronjob.spec</code>	yaml to insert into the created CronJob spec. If specified, this will override any other specified values for the CronJob.	
<code>pingfederate-engine.cronjob.spec.jobTemplate</code>	yaml to override default jobTemplate. If a jobTemplate is not overridden, a default template will be inserted.	
<code>pingfederate-engine.cronjob.image</code>	Image to run the Jobs. The image must include kubectl	<code>bitname/kubectl:latest</code>
<code>pingfederate-engine.cronjob.args</code>	Job arguments	<code>[]</code>
<code>pingfederate-engine.cronjob.podSecurityContext</code>	securityContext for the pod in the jobTemplate. This will be used if a jobTemplate is not specified.	<code>null</code>
<code>pingfederate-engine.cronjob.podSecurityContext</code>	securityContext for the container in the jobTemplate. This will be used if a jobTemplate is not specified.	<code>allowPrivilegeEscalation: false, capabilities: drop: ALL</code>
<code>pingdirectory</code>	PingDirectory values	
<code>pingdirectory.enabled</code>	Enable PingDirectory deployment	<code>false</code>
<code>pingdirectory.cronjob</code>	CronJobs run a kubectl exec command to run commands on a utility sidecar container. They will also create the necessary ServiceAccount, Role, and RoleBinding to run the jobs	
<code>pingdirectory.cronjob.enabled</code>	Enable the PingDirectory CronJob	<code>false</code>
<code>pingdirectory.cronjob.spec</code>	yaml to insert into the created CronJob spec. If specified, this will override any other specified values for the CronJob.	
<code>pingdirectory.cronjob.spec.jobTemplate</code>	yaml to override default jobTemplate. If a jobTemplate is not overridden, a default template will be inserted.	
<code>pingdirectory.cronjob.image</code>	Image to run the Jobs. The image must include kubectl	<code>bitname/kubectl:latest</code>
<code>pingdirectory.cronjob.args</code>	Job arguments	<code>[]</code>
<code>pingdirectory.cronjob.podSecurityContext</code>	securityContext for the pod in the jobTemplate. This will be used if a jobTemplate is not specified.	<code>null</code>

Name	Description	Default
pingdirectory.cronjob.podSecurityContext	securityContext for the container in the jobTemplate. This will be used if a jobTemplate is not specified.	allowPrivilegeEscalation: false, capabilities: drop: ALL
pingdirectory.services.serviceName.loadBalancerService	If true, the per-Pod LoadBalancer services enabled with pingdirectory.services.loadBalancerServicePerPod will include this port.	false
pingdirectory.services.loadBalancerServicePerPod	Set to true to create a separate LoadBalancer service for each individual Pod in the PingDirectory StatefulSet.	false
pingdirectory.services.loadBalancerExternalDNSHostnameSuffix	Value used for the external-dns.alpha.kubernetes.io/hostname annotation for the LoadBalancer services. This value will be used as a suffix for the hostname for each individual pod when pingdirectory.services.loadBalancerServicePerPod is set to true.	
pingdirectoryproxy	PingDirectoryProxy values	
pingdirectoryproxy.enabled	Enable PingDirectoryProxy deployment	false
pingdelegator	PingDelegator values	
pingdelegator.enabled	Enable PingDelegator deployment	false
pingdatasync	PingDataSync values	
pingdatasync.enabled	Enable PingDataSync deployment	false
pingauthorize	PingAuthorize values	
pingauthorize.enabled	Enable PingAuthorize deployment	false
pingauthorizepap	PingAuthorizePAP values	
pingauthorizepap.enabled	Enable PingAuthorizePAP deployment	false
pingaccess-admin	PingAccess admin values	
pingaccess-admin.enabled	Enable PingAccess admin deployment	false
pingaccess-engine	PingAccess engine values	
pingaccess-engine.enabled	Enable PingAccess engine deployment	false

Name	Description	Default
<code>pingcentral</code>	PingCentral values	
<code>pingcentral.enabled</code>	Enable PingCentral deployment	<code>false</code>
<code>pingdataconsole</code>	PingDataConsole values	
<code>pingdataconsole.enabled</code>	Enable PingDataConsole deployment	<code>false</code>
<code>pingdataconsole.defaultLogin</code>	Default login details for the console	
<code>pingdataconsole.defaultLogin.server.host</code>	Default hostname	<code>pingdirectory-cluster</code>
<code>pingdataconsole.defaultLogin.server.port</code>	Default port	<code>636</code>
<code>pingdataconsole.defaultLogin.username</code>	Default username	<code>administrator</code>
<code>PingIntelligence</code>	values	
<code>pingintelligence.enabled</code>	Enable PingIntelligence deployment	<code>false</code>
<code>pd-replication-timing</code>	PingDirectory replication timing values	
<code>pd-replication-timing.enabled</code>	Enable PingDirectory replication timing deployment	<code>false</code>
<code>pingtoolkit</code>	PingToolkit values	
<code>pingtoolkit.enabled</code>	Enable PingToolkit deployment	<code>false</code>
<code>testFramework.rbac.serviceAccountImagePullSecrets</code>	Repository authentication using secrets defined as a docker-registry secrets in Kubernetes.	<code>[]</code>

OpenShift Configuration

Openshift is designed to use a randomly generated user ID and group ID (UID/GID) for the `runAsUser` and `fsGroup` fields of the pod- and container-level security contexts.

By default, the security contexts in the chart use values corresponding to the user and group IDs under which the product runs. You can unset the `fsGroup` and `runAsUser` `securityContext` fields in your custom values, allowing OpenShift to set them as expected.

Unset `fsGroup` and `runAsUser` at the Pod Level

In the global section of the `values.yaml` file, add the following stanza:

```
global:  
  workload:  
    securityContext:  
      fsGroup: null  
      runAsUser: null
```

This will unset `fsGroup` and `runAsUser` in the pod-level security context. Pods that require initContainers will have to also unset `runAsUser` in the container-level security context.

initContainers: Unset runAsUser at the Container Level

Some of the product deployments use initContainers for various operations, such as waiting for other services to be available or configuration actions. These containers, while part of the workload, have the security context set at the container level, not the pod level. The values listed above apply only to the pod-level security context. To unset `runAsUser` for any pingtoolkit initContainers so Openshift can take over, also add the following stanza:

```
global:  
  externalImage:  
    pingtoolkit:  
      securityContext:  
        runAsUser: null
```

For example, here is a complete block for configuring pingaccess-admin with a `waitFor` initContainer:

```
global:  
  workload:  
    securityContext:  
      fsGroup: null  
      runAsUser: null  
  externalImage:  
    pingtoolkit:  
      securityContext:  
        runAsUser: null  
  
pingaccess-admin:  
  enabled: true  
  privateCert:  
  generate: true  
  envs:  
    SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git  
    SERVER_PROFILE_PATH: baseline/pingaccess  
  container:  
    waitFor:  
      pingfederate-engine:  
        service: https  
        timeoutSeconds: 300
```

Examples

Helm Chart Examples

The example deployments of various configurations and combinations of Ping products with these charts are provided on the [Ping Identity DevOps Portal](#).

 **Note**

Install and configure Helm according to the [Getting Started](#) page before trying the examples.

How To

Updating Product Image Tags

This page describes recommendations for updating products and image tags in a Helm installation. The focus is on what is needed to move from one Docker image tag to another in the Helm chart, with some specific steps to follow for certain products. This page does not cover specifics on how to test upgrades or details on how to employ a blue-green update strategy. These processes will depend on the environment where the chart is deployed.

Warning

Always test upgrades in a non-production environment first.

Before updating tags, make sure that you have exported any configuration from your servers and stored it in a server profile. This ensures no configuration is lost when pods are restarted. Information on how to save configuration into a server profile can be found in the [How To](#) section of the [DevOps documentation](#).

The upgrade process depends on the specific product within the Helm chart that is being upgraded.

Blue-Green Update

A simple way to handle version updates is via a blue-green deployment. In this strategy, a full second deployment of the workloads being updated is deployed on the new version, and then traffic is switched from the original deployment to the new one. When following this strategy, be sure that any necessary configuration is captured in a server profile, so that it is maintained in the fresh deployment.

This strategy is particularly suited for stateless applications. Stateful applications like PingDirectory and PingDataSync would require further steps to ensure any data in the current deployment is maintained in the newer deployment.

In-Place Update

Tags and product versions can also be updated in place. The required process depends on the product being updated.

Simple Tag Update

For some products, simply updating the tag is all that is needed. For example, updating from PingDirectory `8.3.0.5-latest` to PingDirectory `9.0.0.0-latest` can be done by updating the tag in values.yaml.

```
pingdirectory:  
  enabled: true  
  image:  
    tag: 8.3.0.5-latest
```

becomes

```
pingdirectory:  
  enabled: true  
  image:  
    tag: 9.0.0.0-latest
```

The active Helm release can be updated after setting the new image tag.

```
helm upgrade --install <releasename> pingidentity/ping-devops -f <updated values.yaml>
```

Update logic will be automatically handled by hook scripts on container startup. This is true for PingDirectory, PingDirectoryProxy, PingDataSync, and PingAuthorize.

For products that are not run as a StatefulSet, the tags can generally be updated without any additional update steps, after the server profile is updated and saved. If the server profile is not up to date with the configuration on the running pod(s), then the configuration could be lost on a tag update.

Product Updates Requiring Manual Steps

Some products have more specific upgrade processes that need to be followed when the product version is being updated, such as updating from PingFederate 10.3.7 to PingFederate 11.0.3.

PingFederate

See the DevOps documentation for [instructions on upgrading PingFederate](#).

By default, `pingfederate-admin` and `pingfederate-engine` each run as Deployments in the chart. If you are running a maintenance upgrade (from PingFederate version x.x.1 to x.x.2 for example), it may be sufficient to export the configuration from the previous version and pull that same configuration via server profile when starting up the new version.



Warning

When upgrading a major or minor version, the upgrade steps at the above link should be followed.

PingAccess

PingAccess follows a similar process to the above instructions for PingFederate.

By default, `pingaccess-admin` runs as a StatefulSet in the chart, so the upgrade utility must be run to update the files in the persistent volume. Once the admin is updated, the engines (which run as Deployments by default) can be restarted with the new tag. See the [PingAccess documentation](#) for more information.

PingAuthorizePAP

Ensure your policies are backed up along with your server profile before updating the tag.

Release Notes

Current Release

Release 0.11.8 (June 4, 2025)

Features

- Updated default global image tag to 2505 .
- Support creation of CronJob resources for PingFederate, similar to the existing support available for PingDirectory (PDI-2218).

Enhancements

- Allow setting securityContext for CronJob and utility sidecar (PDI-2208).

Bug fixes

- Fix incorrect indentation of global configMap name when setting global annotations (PDI-2215).

Previous Releases

Release 0.11.7 (May 1, 2025)

Features

- Updated default global image tag to 2504 .

Release 0.11.6 (April 2, 2025)

Features

- Updated default global image tag to 2503 .

Release 0.11.5 (March 3, 2025)

Features

- Updated default global image tag to 2502 .

Release 0.11.4 (February 14, 2025)

Features

- Updated default global image tag to 2501 . 1 .

Release 0.11.3 (February 3, 2025)

Features

- Updated default global image tag to 2501 .

Bug fixes

- Prevented unexpected service port entries from being created on workloads when setting service labels or annotations. (PDI-2182)

Release 0.11.2 (January 3, 2025)

Features

- Updated default global image tag to 2412 .

Release 0.11.1 (December 3, 2024)

Features

- Updated default global image tag to 2411 .

Release 0.11.0 (November 20, 2024)

Features

- Updated default global image tag to 2410 .

Enhancements

- Added supported values for specifying more fine-grained annotations and labels for workloads, workload pod templates, services, HPA, and RBAC objects. Annotations can now be specified for these resources individually, rather than relying on global annotations that apply to all resources. Here is an example showing the new values that can be set for annotations (analogous values are available to control labels):

```
global:
  annotations:
    globalAnnotation: val

workload:
  annotations:
    globalWorkloadAnnotation: val
  statefulSet:
    annotations:
      globalSSAnnotation: val
  deployment:
    annotations:
      globalDepAnnotation: val

rbac:
  generateGlobalServiceAccount: true
  generateGlobalRoleAndRoleBinding: true
  serviceAccountAnnotations:
    globalServiceAccountAnnotation: val
  roleAnnotations:
    globalRoleAnnotation: val
  roleBindingAnnotations:
    globalRoleBindingAnnotation: val

services:
  annotations:
    globalServiceAnnotation: val
  clusterServiceAnnotations:
    globalClusterServiceAnnotation: val

pingdirectory:
  enabled: true
  annotations:
    pdAnnotation: val
workload:
  annotations:
    pdWorkloadAnnotation: val
  statefulSet:
    annotations:
      pdSSAnnotation: val
rbac:
  generateServiceAccount: true
  generateRoleAndRoleBinding: true
  role:
    rules:
      - apiGroups: []
        resources: ["secrets"]
        verbs: ["get", "list"]
  serviceAccountAnnotations:
    pdServiceAccountAnnotation: val
  roleAnnotations:
    pdRoleAnnotation: val
  roleBindingAnnotations:
    pdRoleBindingAnnotation: val
services:
  annotations:
    pdServiceAnnotation: val
  clusterServiceAnnotations:
    pdClusterServiceAnnotation: val
```

Bug fixes

- Fixed global annotations not applying for RBAC objects.

Release 0.10.9 (October 1, 2024)

Features

- Updated default global image tag to 2409 .

Release 0.10.8 (September 4, 2024)

Features

- Updated default global image tag to 2408 .

Release 0.10.7 (August 6, 2024)

Features

- Updated default global image tag to 2407 .

Release 0.10.6 (July 2, 2024)

Features

- Updated default global image tag to 2406 .

Release 0.10.5 (June 5, 2024)

Features

- Updated default global image tag to 2405 .

Release 0.10.4 (May 1, 2024)

Features

- Updated default global image tag to 2404 .

Release 0.10.3 (March 29, 2024)

Features

- Updated default global image tag to 2403 .

Release 0.10.2 (March 1, 2024)

Features

- Updated default global image tag to `2402`.

Release 0.10.1 (February 5, 2024)

Bug Fixes

- Fixed templating failure when not specifying `ingress.spec.ingressClassName` with Ingress enabled.

Release 0.10.0 (January 31, 2024)

Features

- Updated default global image tag to `2401`.
- Added support for setting environment variables in utility sidecar pods, with the `utilitySidecar.env` value.
- Added support for setting `ingressClassName` in Ingress specs with the `ingress.spec.ingressClassName` value.

Enhancements

- Updated generated PVC definitions to include annotations.

Release 0.9.22 (December 29, 2023)

Features

- Updated default global image tag to `2312`.

Release 0.9.21 (December 4, 2023)

Features

- Updated default global image tag to `2311`.

Defects

- Updated the workload template to avoid setting `replicas` when autoscaling is enabled.
- Improved capabilities checks for `apiVersion` field to avoid issues with prerelease Kubernetes versions.

Release 0.9.20 (November 2, 2023)

Features

- Updated default global image tag to `2310`.

-
- Added environment variables for PingDirectoryProxy to support enabling automatic server discovery.

Defects

- Updated the CronJob template to handle the switch from `batch/v1beta1` to `batch/v1` in Kubernetes 1.25.

Release 0.9.19 (September 6, 2023)

Features

- Updated default global image tag to `2308`.

Release 0.9.18 (August 28, 2023)

Resolved Defects

- Fixed incorrect yaml formatting when setting `testFramework.rbac.serviceAccountImagePullSecrets`.

Release 0.9.17 (August 25, 2023)

Features

- Added support for setting `imagePullSecrets` in workloads.
- Added support for setting `testFramework.rbac.serviceAccountImagePullSecrets` to add secrets to the testFramework service account.

Release 0.9.16 (August 2, 2023)

Features

- Updated default global image tag to `2307`.

Release 0.9.15 (July 13, 2023)

Features

- Updated default global image tag to `2306`.

Enhancements

- Updated template to allow setting a custom workload type when using a HorizontalPodAutoscaler.

Release 0.9.14 (June 2, 2023)

Features

- Updated default global image tag to `2305`.

Release 0.9.13 (May 4, 2023)

Features

- Updated default global image tag to `2304`.

Release 0.9.12 (April 3, 2023)

Features

- Updated default global image tag to `2303`.
- Updated the workload `topologySpreadConstraints` field to automatically set `matchLabels` to match the workload labels.

Release 0.9.11 (March 3, 2023)

Features

- Updated default global image tag to `2302`.
- Added default environment variables to `pingdirectoryproxy` to support joining a PingDirectory topology.

Release 0.9.10 (February 3, 2023)

Features

- Updated default global image tag to `2301`.
- Updated the securityContext defaults for Pods and containers in the ping-devops Helm chart to satisfy the "restricted" Pod Security Standard in Kubernetes.
- Added support for running a separate LoadBalancer service for each PingDirectory pod. This may be useful when running across multiple regions when using VPC peering isn't possible.

Resolved Defects

- Updated the HorizontalPodAutoscaler API to use the correct value for Kubernetes versions greater than 1.23.

Release 0.9.9 (January 3, 2023)

Features

- Updated default global image tag to `2212`.
- Removed `pingdatagovernance` and `pingdatagovernancepap` from the chart. Use `pingauthorize` and `pingauthorizepap` instead.

Release 0.9.8 (December 5, 2022)

Features

- Updated default global image tag to `2211`.
- Custom annotations can now be specified for Services.

Defects

- Fixed HorizontalPodAutoscaler autoscalingMetricsTemplate being inserted in the wrong location in the generated yaml.
- Fixed the documentation in values.yaml referring to `pingdirectory.cronjob.jobspec` rather than the correct value `pingdirectory.cronjob.jobTemplate`.

Release 0.9.7 (November 2, 2022)

Features

- Updated default global image tag to `2210`.

Release 0.9.6 (October 4, 2022)

Features

- Updated default global image tag to `2209`.
- Added support for deploying a HorizontalPodAutoscaler for pingaccess-engine, pingfederate-admin, pingdelegator, pingauthorize, pingauthorizepap, pingcentral, and pingdataconsole. Previously, deploying a HorizontalPodAutoscaler was only supported for pingfederate-engine.
- Added support for setting the podManagementPolicy for StatefulSet workloads. The default policy is OrderedReady. The Parallel policy allows for starting up multiple Pods of the StatefulSet simultaneously, improving initial deployment times. Parallel startup is only supported with PingDirectory and PingDataSync, and only with images version 2209 and newer.

Release 0.9.5 (September 1, 2022)

Features

- Updated default global image tag to `2208`.
- Added support for setting container-level securityContext values for the main container of each workload. By default no container-level securityContext will be set. A container-level securityContext isn't necessary if the values from the Pod-level securityContext are sufficient.
- Added support for setting the topologySpreadConstraints field on workloads.
- Added support for setting the enableServiceLinks field on workloads.

Documentation

- Added an [example](#) for mounting keystore secrets with Vault.
- Added an [example](#) for mounting secrets with CSI volumes (which can be used for various storage systems including AWS secrets manager)
- Fixed Helm [RBAC example](#) using an invalid serviceAccountName for pingauthorize.
- Added a [doc page](#) describing how to update product versions.
- Added example docs for deploying [PingDirectory](#) and [PingFederate](#) in a multi-region environment with Helm.

Resolved Defects

- Removed support for apache-jmeter, since it is better suited to run as a job than as a long-running workload.

Release 0.9.4 (August 5, 2022)

Features

- Updated default global image tag to [2207](#).
- Added support for apache-jmeter

Resolved Defects

- Fixed an issue making it impossible to use an existing service account (an account not managed by the Helm chart) for a workload. An existing service account can now be used by specifying the {product}.rbac.serviceAccountName field while leaving {product}.rbac.generateServiceAccount set to the default false value. See the PingAuthorize section of the updated [RBAC example](#).

Release 0.9.3 (July 1, 2022)

Features

- Updated default global image tag to [2206](#).
- Added support for PingIntelligence.
- Updated the Helm chart to support generating ServiceAccounts, Roles, and RoleBindings for a workload. These can be generated globally (one common to each workload) or individually for each workload. By default, none will be generated.

These can be controlled with the global.rbac (or {product}.rbac) section. To generate a common ServiceAccount usable by all workloads, use `global`.rbac.generateGlobalServiceAccount``. Use `rbac.generateServiceAccount` to generate separate ServiceAccounts for individual workloads. Similarly, use `global.rbac.generateGlobalRoleAndRoleBinding` and `rbac.generateRoleAndRoleBinding` for creating a Role and RoleBinding.

Set `rbac.applyServiceAccountToWorkload` to true to set the account on the Deployment or StatefulSet. The name of the ServiceAccount will be autogenerated unless the `rbac.serviceAccountName` field is set. The specific Role yaml can be provided in `rbac.role`.

The Vault default has changed. The Vault serviceAccount will now default to the autogenerated account for the workload, instead of the previous default of "vault-auth". This can be overriden by setting the `vault.hashicorp.annotations.serviceAccountName` value.

See the table with sample Helm chart value files found on the [Ping Identity Devops Portal](#) for an example.

- Added a default empty `global.labels` section.

Release 0.9.2 (June 2, 2022)

Features

- Updated default global image tag to `2205`.
- Added support for providing a null securityContext for a workload, which is useful for OpenShift security context constraints.
- Added support for enabling Ingress for `pingdirectoryproxy` and `pingdatasync`.
- Updated `pingdirectoryproxy` to be a StatefulSet by default in the Helm charts, with the persistent volume disabled. This supports having consistent proxy pod names.
- Added a new `privateCert.format` field, which can be set to "pingaccess-fips-pem" to generate a cert that can be used by PingAccess when running in FIPS mode. The cert is generated by a temporary PingAccess initContainer, as PingAccess requires a specific format for certs when in FIPS mode that must be generated from PingAccess itself. Leaving the field blank or setting it to any other value will generate a cert in the same manner as before: adding the key pair to a PKCS12 keystore file.
- Updated the `externallImage` values section to expect the same format for `image:` as the individual products (repository, image name, tag, etc. are provided separately). If no image values are specified for an `externallImage`, the corresponding defaults from the main product section will be used. For example, if `global.externalImage.pingtoolkit.image` is empty, then the values from the top-level `pingtoolkit.image` section will be used.

Release 0.9.1 (May 5, 2022)

Features

- Updated default global image tag to `2204`.
- Updated the `PingDataSync` env vars ConfigMap to include variables needed to enable failover between servers. Failover will be enabled when deploying two or more `PingDataSync` replicas
- Reduced utilitySidecar resource requests.

Resolved Defects

- Updated the `image.repositoryFqn` field to be consistent with the other fields under `image`. Previously, `repositoryFqn` was expected at the same level as the `image:` section, now it is expected within the `image:` section like other fields (tag, pullPolicy, etc.). The image tag must now be provided separately from the `repositoryFqn`. The `repositoryFqn` should only be the name of the repository, not the tag of the specific image.
- Fixed a version check in the Helm chart for choosing the correct k8s API for Ingress. The version check was previously failing on EKS clusters due to the format EKS uses for the cluster version.

Release 0.9.0 (April 1, 2022)

Features

- Default global image tag updated to 2203 .
- Customizability on Cronjob and Utility Sidecar:
 - Override jobTemplate in CronJob now available.
 - Override image used in utilitySidecar now available.
- Updated the default PingDataSync workload in the Ping devops Helm charts to use a StatefulSet rather than a Deployment. This ensures that the sync-state.ldif file is maintained between pod restarts.

Release 0.8.9 (March 17, 2022)

Features

- Edit from 0.8.8 release. Previously, the image fully qualified name also included the image tag, which was then duplicated upon deployment when "tag" value present.

Release 0.8.8 (March 16, 2022)

Features

- Added support for fully qualified image location. For more information go to the image section in our [values.yaml](#).

```
image:
  repository: pingidentity
  repositoryFqn:
  name:
  tag: "2202"
  pullPolicy: IfNotPresent
```

Release 0.8.7 (March 11, 2022)

Features

- Corrected default global image tag updated to 2202 .

Release 0.8.6 (March 3, 2022)

Features

- Default global image tag updated to 2202 .

Release 0.8.5 (February 7, 2022)

Features

- PingCentral now supported. Example values application found [here](#).

Issues Resolved

- [Issue #119](#) — Workload template not honoring false values from values.yaml. Previously, false did not overwrite true in the Ping Identity Helm Chart template. This fix in _merge-util.tpl will resolve multiple cases within the Ping Identity Helm Chart.

```
{{- $globalValues := deepCopy $top.Values.global -}}
{{- $prodValues := deepCopy (index $top.Values $prodName) -}}
{{- $mergedValues := mergeOverwrite $globalValues $prodValues -}}
```

- [Issue #264](#) — Update default global.image.tag to 2201.

Release 0.8.4 (January 7, 2022)

Issues Resolved

- Fixed an issue that caused installation to fail when enabling pingtoolkit.

Release 0.8.3 (January 6, 2022)

Features

- Document [supported values](#).

Issues Resolved

- [Issue #233](#) — Ingress - semverCompare now retrieves correct K8 version for applying the correct apiVersion.

```
{{- if semverCompare ">=1.19.x" $top.Capabilities.KubeVersion.Version }}
```

- [Issue #254](#) — Update default global.image.tag to 2112.

Release 0.8.2 (December 17, 2021)

- [Issue #238](#) — Added support for running a utility sidecar alongside a product workload

The `utilitySidecar` field under a given product can be used to run a sidecar container that will permanently alongside the product container. This sidecar can be used for utility command-line processes, such as running the collect-support-data tool or running a backup.

An example can be found in the docs/examples/pingdirectory-backup directory for running a PingDirectory backup every 6 hours via a CronJob.

```
pingdirectory:  
  workload:  
    shareProcessNamespace: true  
  utilitySidecar:  
    enabled: true
```

- [Issue #247](#) — Update default global.image.tag to 2111.1.

Release 0.8.1 (December 6, 2021)

- [Issue #240](#) — Fixed failure on installation of 0.8.0 due to missing PingDirectory HTTP port value.

Release 0.8.0 (December 6, 2021)

- [Issue #229](#) — Support for `shareProcessNamespace` in pod spec.

A PingDirectory utility sidcar container needs to share the process namespace with the main PingDirectory container running in the same pod in order to get useful output out of tools like jps. More support to come on the utility sidcar in future Helm release.

- [Issue #232](#) — Update default global.image.tag to 2111
- [Issue #239](#) — Support for custom container arguments

```
pingfederate-admin:  
  enabled: true  
  container:  
    args: ["start-server", "tail -f /dev/null"]
```

- [Issue #232](#) — Update default global.image.tag to 2111
- [Issue #240](#) — Allow specifying PingDirectory HTTPS port in values

```
pingdirectory:  
  enabled: true  
  services:  
    https:  
      containerPort: 8443
```

Release 0.7.9 (December 1, 2021)

- [Issue #223](#) — Support for HPA Scaling Behavior

```
clustering:  
  autoscaling:  
    enabled: true  
    behavior:  
      scaleDown:  
        stabilizationWindowSeconds: 300  
        policies:  
          - type: Percent  
            value: 100  
            periodSeconds: 15  
      scaleUp:  
        stabilizationWindowSeconds: 0  
        policies:  
          - type: Percent  
            value: 100  
            periodSeconds: 15  
          - type: Pods  
            value: 4  
            periodSeconds: 15  
        selectPolicy: Max
```

- [Issue #231](#) — Helm test image pull policy no longer hard-coded in `helm-charts/charts/ping-devops/templates/pinglib/_tests/tpl`.

```
- imagePullPolicy: IfNotPresent
```

- [Issue #233](#) — Cluster service for pingaccess-admin in Multi-region Support for multi-region PingAccess deployment without using an ingress. The headless service is an effective way to share the pod id across clusters.

```
pingaccess-admin:
  enabled: true
  privateCert:
    generate: true
  envs:
    SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git
    SERVER_PROFILE_PATH: baseline/pingaccess
  container:
    replicaCount: 1
    waitFor:
      pingfederate-engine:
        service: https
  services:
    https:
      servicePort: 9000
      containerPort: 9000
      ingressPort: 443
      dataService: true
      clusterService: true
    clusterconfig:
      servicePort: 9090
      containerPort: 9090
      ingressPort: 443
      dataService: true
  clusterExternalDNSHostname: pingaccess-admin.usa.ping-multi-cluster.com
```

Release 0.7.8 (November 2, 2021)

- [Issue #213](#) — Removed default SERVER_PROFILE variables from values.yaml

```
envs:
- SERVER_PROFILE_URL:
- SERVER_PROFILE_PATH:
```

- [Issue #216](#) — Add option to generate a master password for ping services

In the interest of better security practice, this enhancement provides the ability to generate this password via the derivedPassword function in helm. With this, several items can be used by default and overridden by the deployer to generate a secure password. When it generates the password:

- A note will be added to the NOTES (see below)
- The password will be set into the global configmap PING_IDENTITY_PASSWORD. (we may want to use a secret instead)

NOTES (see the generated password as well as the WARNING)

```
$ helm upgrade --install test-pw pingidentity/ping-devops --set global.masterPassword.enabled=true
NOTES:
#-----
# Ping DevOps
#
# Description: Ping Identity helm charts - 09/18/21
#
# WARNING: Master Password has been requested and generated. This is intended to
#           generate a password for DEVELOPMENT PURPOSES ONLY. This password will be
#           assigned to the PING_IDENTITY_PASSWORD unless overridden by the values.
#
#           PING_IDENTITY_PASSWORD: *****
#-----
```

The values used to drive the creation of this password are:

values.yaml

```
global:
#####
# Master Password Generation
#
# Uses Helm function derivePassword, which uses the master password
# specification: https://masterpassword.app/masterpassword-algorithm.pdf
#
#   masterPassword.enabled: {true | false}
#   masterPassword.strength: {master password template: long | maximum}
#   masterPassword.name: {defaults to .Release.Name}
#   masterPassword.site: {defaults to .Chart.Name}
#   masterPassword.secret: {defaults to .Release.Namespace}
#####
masterPassword:
  enabled: false
  strength: long
  name: # default - .Release.Name
  site: # default - .Chart.Name
  secret: # default - .Release.Namespace
```

+ As shown in the example above, a deployer only needs to provide the global.masterPassword.enabled=true to have it generated.

- [Issue #221](#) — PingDirectory service.x.containerPort updates to LDAPS_PORT environment variable.
- [Issue #222](#) — Update default global.image.tag to 2110 .
- [Issue #224](#) — External Hostname Annotations on PD data service . == Release 0.7.7 (Oct 7, 2021)
- [Issue #217](#) — Update default security context group id to root (0).

```
global:  
  workload:  
    securityContext:  
      fsGroup: 0  
      runAsUser: 9031  
      runAsGroup: 0
```

- [Issue #218](#) — Update default global.image.tag to 2109 .

Release 0.7.6 (September 18, 2021)

- [Issue #209](#) — Fix incorrect default ldap-sdk-tools probe exec commands.
- [Issue #210](#) — Add helm-chart product/image pingtoolkit.
- [Issue #211](#) — Allow for schedulerName to be provide on workloads (pods).

Release 0.7.5 (August 30, 2021)

- [Issue #206](#) — Bump default image tag to 2108 .

Release 0.7.4 (August 26, 2021)

- [Issue #196](#) — Set initContainer settings from values.yaml instead of hard coded templates.

This issue was created since the initContainer resources were hard coded in the template, not allowing the implementor to provide their own values, causing issues when trying to deploy the pingfederate-engine in openshift.

Moving a lot of the hard coded yaml out of the template files into the default values.yaml file. This will give the implementor full control of how the initContainer runs.

One breaking change with the values.yaml if anyone has overridden, is that the `{image name}` in the `global.externalImage.{name}: {image name}` value is moved into a map. The default pingtoolkit externalImage looks like:

```
global:  
  externalImage:  
    pingtoolkit:  
      image: pingidentity/pingtoolkit:2107  
      imagePullPolicy: IfNotPresent  
      resources:  
        limits:  
          cpu: 1m  
          memory: 128Mi  
        requests:  
          cpu: 500m  
          memory: 64Mi  
      securityContext:  
        allowPrivilegeEscalation: false  
        capabilities:  
          drop:  
            - ALL  
      readOnlyRootFilesystem: true  
      runAsNonRoot: true  
      runAsUser: 9031  
      runAsGroup: 9999
```

- [Issue #203](#) — testFramework - Support multiple waitFor products in testSteps.

When there are two waitFor's together, allow for combining them to run them within same initContainer, with a definition like:

```
testSteps:  
  - name: 01-wait-for  
    waitFor:  
      pingfederate-admin:  
        service: https  
      pingfederate-engine:  
        service: https
```

creating a couple of initContainers of:

```
initContainers:  
  - name: 01-wait-for-pingfederate-admin  
    ...  
  - name: 01-wait-for-pingfederate-engine  
    ...
```

Release 0.7.3 (August 24, 2021)

- [Issue #194](#) — Change default envs for pingauthorize/pingauthorizepap.

The current envs for pingauthroize in the values.yaml file are:

```
envs:  
  SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git  
  SERVER_PROFILE_PATH: paz-pap-integration/pingauthorize  
  SERVER_PROFILE_PARENT: PAZ  
  SERVER_PROFILE_PAZ_URL: https://github.com/pingidentity/pingidentity-server-profiles.git  
  SERVER_PROFILE_PAZ_PATH: baseline/pingauthorize
```

Just a side note here, the `baseline/pingauthorize` PATH includes a connection to pingdirectory, which will cause this to fail (pingauthorize https will return a 503).

If someone wants to override these, they need to be sure to use/override the SERVER_PROFILE_PARENT variable, so the parent profiles aren't brought in.

The better default values.yaml should probably be:

```
envs:  
  SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git  
  SERVER_PROFILE_PATH: getting-started/pingauthorize
```

For pingauthorizepap, it should have a default SERVER_PROFILE variables as empty, as no SERVER_PROFILE is needed by default.

- [Issue #198](#) — testFramework: Support full definition of initContainers attributes in testSteps and finalStep.

Update the testFramework to pull in all attributes of the testSteps and finalStep into the init containers and final container. This allows for setting any resource, imagePullPolicy, ...

This came about as there was no way to set resource or imagePullPolicy details.

With this change, will be adding a couple of defaults into the value.yaml file for the finalStep:

```
finalStep:  
  name: 99-completion  
  image: busybox  
  imagePullPolicy: IfNotPresent  
  command:  
    ...  
  resources:  
    limits:  
      cpu: 500m  
      memory: 128Mi  
    requests:  
      cpu: 1m  
      memory: 64Mi
```

Release 0.7.2 (August 13, 2021)

- [Issue #191](#) — Change variable PF_ADMIN_BASEURL to PF_ADMIN_PUBLIC_BASEURL.

Release 0.7.2 created the new variable `PF_ADMIN_BASEURL`. Due to the current user of the same variable with added `_PUBLIC_`, the actual variable name needs to be `PF_ADMIN_PUBLIC_BASEURL`.

Release 0.7.1 (August 13, 2021)

- [Issue #187](#) — Create the PUBLIC hostname/ports in the global env vars configmap all the time.

Currently, the PUBLIC hostname/ports in the global env vars configmap are created if and only if the ingress is enabled.

Normally, this would be fine, except that some of the products (i.e PingFederate) use the PUBLIC environment variable to setup items like BASE URLs and redirects for the browser. This is required for use cases when there is no ingress, but the user creates a port forward, as well as testing with no ingresses.

So, if no ingress is created, then the PUBLIC_HOSTNAMES should be set to localhost and the PUBLIC_PORT_* should be set to the same port as the containerPort.

If ingress is used, then the functionality will not be changed, and the public hostname will be constructed as well as the public ingressPort.

- [Issue #188](#) — Add the PF_ADMIN_BASEURL environment variable to the pingfederate admin/engine configmaps

With the 10.3 release of PingFederate, there is a variable used to provide redirect links called the PF_ADMIN_BASEURL. This needs to be set by the helm chart, as it will either be a public host or localhost, depending on if the ingress is available. The container has no idea which it should be as it doesn't have insight into the environment it's running.

If ingress is enabled, an example for this variable is:

```
PF_ADMIN_BAESURL=https://pingfederate-admin.example.com
```

If ingress is not enabled, an example for this variable:

```
PF_ADMIN_BASEURL=https://localhost:9999
```

Release 0.7.0 (August 9, 2021)

- [Issue #184](#) — Create default ServiceAccount/Role/RoleBinding for testFramework.

To allow for a role to be created during testing, an rbac section is added to the testFramework allowing for the definition of that Role. If enabled, it will create a ServiceAccount, Role and RoleBinding using the same naming rules of resources and add that serviceAccount to the test pod.

testFramework default rbac set to:

```
#####
# If rbac is enabled, this will create:
#   - serviceAccount
#   - role
#   - roleBinding (between serviceAccount and role)
#
# and apply the serviceAccount to the pod in the tests.
# The names for these resources will be named using the
# naming rules for all resources including the ReleaseName
#####
rbac:
  enabled: true
  role:
    rules:
      - apiGroups:
          - '*'
        resources:
          - '*'
        verbs:
          - '*'
```

Release 0.6.9 (August 6, 2021)

- [Issue #179](#) — Bump default image tag to 2107 Issue #182 Set default startupProbe.timeoutSeconds to 5.
- [Issue #180](#) — Enhance testFramework to support additional pod level configurations.

When using the testFramework there are additional pod level config items that need to be provided (i.e. serviceAccountName) along with the existing securityContext.

To allow for any item to be configured, we should add a testFramework.pod that will pull in all items into the testFramework pod definition.

Example:

```
testFramework:
#####
# Pod information to include
#
# Examples:
#   securityContext for all containers
#   serviceAccount for all containers
#####
pod:
  securityContext:
    runAsUser: 1000
    runAsGroup: 2000
  serviceAccount: serviceaccount-name
```

Warning

This will be a breaking change for anyone who has created a testFramework.securityContext. If this is the case, they need to add pod in front of securityContext.

Release 0.6.8 (July 29, 2021)

- [Issue #175](#) — Invalid ingress resources on Kubernetes clusters > 1.18.

During resolution of issue #170 providing support for ingress apiVersion v1, the necessary ingress yaml fields weren't updated to reflect that new version. This is a fix. The backend definition of the Ingress will now reflect the proper definition based on a v1 or v1beta1 apiVersion.

Example: If KubeVersion > 1.18

```
service:  
  name: https  
  port:  
    number: 443
```

Example: If KubeVersion ≤ 1.18

```
serviceName: https  
servicePort: 443
```

Additionally, adding the pathType for all versions as it is now required in ingress v1.

Release 0.6.7 (July 28, 2021)

- [Issue #170](#) — Update Ingress resource kind.

If kubernetes version is >1.18, setting the ingress apiVersion to `v1`. Otherwise, current default will be used `v1beta1`.

- [Issue #171](#) — Reevaluate Lifecycle probes.

Adding startupProbe as well as re-organizing how the probes are defined, allowing the deployer to use standard k8s probe definitions out of the box.

- Moving the probes section under global.container
- Changing names: (liveness → livenessProbe, readiness → readinessProbe)
- Adding startupProbe

The new default looks like:

```
#####
# Probes
#
# Probes have a number of fields that you can use to more precisely control the
# behavior of liveness and readiness checks.
#
# https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-startup-probes/
#####
probes:
  livenessProbe:
    exec:
      command:
        - /opt/liveness.sh
    initialDelaySeconds: 30
    periodSeconds: 30
    timeoutSeconds: 5
    successThreshold: 1
    failureThreshold: 4
  readinessProbe:
    exec:
      command:
        - /opt/readiness.sh
    initialDelaySeconds: 30
    periodSeconds: 5
    timeoutSeconds: 5
    successThreshold: 1
    failureThreshold: 4
  startupProbe:
    exec:
      command:
        - /opt/liveness.sh
    periodSeconds: 10
    failureThreshold: 90
```

Breaking Changes

This is a breaking change if anyone has overriding probes in their own values file. The fix is simply move their definition of their probes to live under global.container or the (productName).container, as well as adding "Probe" to the definition.

Release 0.6.6 (July 7, 2021)

- [Issue #160](#) — Change default image tag to 2106 .
- [Issue #166](#) — Add securityContexts to testFramework containers.
 - Adding ability to provide a securityContext at the following levels:
 - Changing the default finalStep image to busybox

```
testFramework:  
...  
#####  
# SecurityContext for all containers  
#####  
securityContext:  
  runAsUser: 1000  
  runAsGroup: 2000  
...  
testSteps:  
- name: 01-init-example  
...  
  securityContext:  
    runAsUser: ...  
...  
finalStep:  
  securityContext:  
    runAsUser: ...
```

- [Issue #167](#) — Disable testFramework by default. To enable, simply:

```
testFramework:  
  enabled: true  
...  
...
```

Release 0.6.5 (July 4, 2021)

- [Issue #163](#) — Add PingAuthorize and PingAuthorizePAP to helm charts.

This includes the pre-release to PingAuthorize 8.3. It includes the necessary config for PingAuthorize and PingAuthorizePAP, even though there isn't a release for 2105. The current edge release is required to use the default server-profiles provided in the values.yaml. Once the global tag is changed to 2106 (over next few days) PingAuthorize will be default for use over PingDataGovernance. This will be tracked in a ticket released 2105.

Example yaml to test PingAuthoize/PAP:

```
pingdataconsole:  
  enabled: true  
  
pingdirectory:  
  enabled: true  
  
pingauthorize:  
  image:  
    tag: 8.3.0.0-edge  
  enabled: true  
  
pingauthorizepap:  
  enabled: true
```

Release 0.6.4 (July 1, 2021)

- [Issue #158](#) — Increment default tag to 2105 Sidecars and initContainers are valuable for a multitude of reasons - log forwarding, metric exporting, backup jobs. Because of this they can also have many ways of being configured.

Allow for defining three top level maps to provide details for:

- sidecars — Defines sidecar containers to be run alongside product containers.
- initContainers — Defines initContainers to be run before product containers.
- volumes — Defines volumes used by sidecars, initContainers and product containers.

Example definitions:

```
sidecars:  
  pd-access-logger:  
    name: pd-access-log-container  
    image: pingidentity/pingtoolkit:2105  
    volumeMounts:  
      - mountPath: /tmp/pd-access-logs/  
        name: pd-access-logs  
        readOnly: false  
  statsd-exporter:  
    name: statsd-exporter  
    image: prom/statsd-exporter:v0.14.1  
    args:  
      - "--statsd.mapping-config=/tmp/mapping/statsd-mapping.yml"  
      - "--statsd.listen-udp=:8125"  
      - "--web.listen-address=:9102"  
    ports:  
      - containerPort: 9102  
        protocol: TCP  
      - containerPort: 8125  
        protocol: UDP  
  
initContainers:  
  init-1:  
    name: 01-init  
    image: pingidentity/pingtoolkit:2105  
    command: ['sh', '-c', 'echo "Initing 1" && touch /tmp/pd-access-logs/init-1']  
    volumeMounts:  
      - mountPath: /tmp/pd-access-logs/  
        name: pd-access-logs  
        readOnly: false  
  
volumes:  
  pd-access-logs:  
    emptyDir: {}  
  statsd-mapping:  
    configMap:  
      name: statsd-config  
      items:  
        - key: config  
          path: statsd-mapping.yml
```

And within the product (or global) definition, allow for inclusion of sidecars, initContainers and volumes. These must be available in the top-level sidecars:, initContainers:, and volumes:

- includeSidecars
- includeInitContainers — Run in order as listed in array
- includeVolumes

Example usages:

```
pingdirectory:  
...  
includeSidecars:  
  - pd-access-logger  
includeInitContainers:  
  - init-1  
includeVolumes:  
  - pd-access-logs  
  
volumeMounts:  
  - mountPath: /opt/access-logs/  
    name: pd-access-logs
```

Release 0.6.3 (June 21, 2021)

- [Issue #154](#) — Increment default tag to 2105 .
- [Issue #155](#) — Add clusterServiceName to product services with service clusters.

Release 0.6.2 (May 24, 2021)

- [Issue #151](#) — Add support for Container LifeCycle Event Hooks.

Adding the following to values.yaml:

```
global:  
  #####  
  # container life handlers, allowing for lifecycle events such  
  # as postStart and preStop events  
  #  
  # https://kubernetes.io/docs/tasks/configure-pod-container/attach-handler-lifecycle-event  
  #####  
  lifecycle: {}  
  # Example  
  # lifecycle:  
  #   postStart:  
  #     exec:  
  #       command: ["/bin/sh", "-c", "echo Start Complete > /tmp/message"]
```

- General cleanup of values.yaml comments.

- Setting default `externalImages.pingtoolkit` tag to `2104`, and removing `edge` tag from `ldap-sdk-tools` image which will now default to same `global.image.tag` setting (currently `2104`).

Release 0.6.1 (May 21, 2021)

- [Issue #148](#) — Calculate checksum of `ConfigMaps` based on the data rather than entire `ConfigMap` file.

This will only use the `ConfigMap.data` when creating checksums in workload rather than using the entire file. It will result in no checksum change when labels/annotations are the only thing changing. A good example is the helm chart version, which changes the label, but not data.

Release 0.6.0 (May 11, 2021)

- Changed default `global.image.pullPolicy` from `Always` to `IfNotPresent`.

This is due to the fact that the `global.image.tag` is a non-floating tag. Once it is downloaded and present, it will not change. This small change will increase performance at startup as images are typically present when installing/updating releases.

Simply set `global.image.pullPolicy=Always` to pull every time if needed.

- BETA 2 - Testing Framework supporting helm test command and associated `testFramework` values.

Cleaned up the generation of resources honoring the `addReleaseNameToResource` setting.

Release 0.5.9 (May 10, 2021)

- BETA 1 - Testing Framework supporting helm test command and associated `testFramework` values.

A testing framework is being created to allow for testing Ping Identity helm chart deployments using a `testFramework` set of values. This is currently in beta, with documentation to available soon. Expect that changes will be made to this work, until it's fully released with documentation.

Release 0.5.8 (May 6, 2021)

- [Issue #141](#) — Fix `DNS_QUERY_LOCATION` on pingfederate-engine configmap.yaml

Resolves an issue with the `DNS_QUERY_LOCATION` when pingfederate clustering is used for >1 pingfederate-engines

Release 0.5.7 (May 3, 2021)

- [Issue #136](#) — ClusterIP Services port/targetPort be set to the containerPort

Since the ClusterIP Services (aka Headless services) only provide access to the underlying container IP and port. The port, and by default `targetPort`, will be set to the `containerPort` value. The helm charts will start requiring the `containerPort` for any service where `clusterService:true` is set, otherwise it will fail with an error message.

- [Issue #138](#) — Update `image.tag` to 2104 (April 2021)

Release 0.5.5 (April 29, 2021)

- [Issue #133](#) — Change default pingdirectory values (container.resources.requests.cpu=50m and container.replicaCount=1)

Setting the cpu request to 50m, will provide at last some reservation of CPU, so that if there are multiple nodes, it will better even out the load.

Additionally, setting the `replicaCount` to 1 by default, as many cases in development, there isn't a great need to have multiple replicas. If this is the case, simply set `pingdirectory.container.replicaCount=2` or any number of replica's.

- [Issue #132](#) - Adding PingDirectoryProxy to mix of products

Release 0.5.5

- [Issue #126](#) — Unable to mount `secretVolume` and `configMapVolumes` simultaneously

This is one additional fix to the the same thing fixed in 0.5.4. `volumeMounts:` had the same issue as `volumes:`. This completes and resolves issue #126.

Release 0.5.4

- [Issue #126](#) — Unable to mount `secretVolume` and `configMapVolumes` simultaneously

Due to the fact that `volumes:` is an array of items, `volumes:` usage with `secret` or `configMap` volumes exposed the issue that multiple `volumes:` entries were used, and only kept the last one. Fix included only using `volumes:` once. Note that the template will end up with a `volumes: null` if none are set (i.e. deployment with no Secret/ConfigMap volumes), but that is ok.

Release 0.5.3

- [Issue #121](#) — Create `global-env-vars` hosts/ports for all products regardless if enabled

The status of this config map is used to form the checksum for the products. This will ensure that a simple addition/deletion of a product from the deployed mix won't cause all products to be restarted.

- [Issue #122](#) - Update `image.tag` to 2103 (March 2021)

The image tag is modified to 2103. This includes:

- Security Context on StatefulSets to include a `fsGroup=9999` (same as `gid`)
- Update the services ContainerPort to unprivileged ports (i.e. 636 → 1636)

Release 0.5.2

- [Issue #113](#) — Default pingaccess-admin to StatefulSet

In order to provide HA with a PingAccess cluster between admin/engine nodes, it is required that the PingAccess Admin deploy as a StatefulSet with persistence. Otherwise if the PingAccess Admin goes down, the engines would lose connectivity to that node and be unable to get further config updates and subsequently have to bounce and lose their web-session information.

The new default yaml

```
pingaccess-admin:  
  workload:  
    type: StatefulSet
```

- [Issue #95](#) — Fix default serviceAccount in workload for vault

Fixed issue that was created in Issue 95 (using annotations to provide vault details) to pull serviceAccountName from the proper location in annotations.

```
vault:  
  hashicorp:  
    annotations:  
      serviceAccountName: vault-auth
```

- [Issue #116](#) — Support Annotations at Workload Level.

Support annotations at the workload level. For workloads, adding `.spec.template.metadata`.

Example telegraf annotation:

```
pingfederate-engine:  
  workload:  
    annotations:  
      telegraf.influxdata.com/class: app
```

would lead to:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/instance: samir
    app.kubernetes.io/managed-by: Helm
    app.kubernetes.io/name: pingfederate-engine
    helm.sh/chart: ping-devops-0.5.1
  name: samir-pingfederate-engine
spec:
  replicas: 1
  selector:
    matchLabels:
      app.kubernetes.io/instance: samir
      app.kubernetes.io/name: pingfederate-engine
  strategy:
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
      type: RollingUpdate
  template:
    metadata:
      annotations:
        telegraf.influxdata.com/class: app
```

- [Issue #117](#) — Bug - cluster service shouldn't use image name for service name.
- [Issue #114](#) — Revamp vault.hashicorp.secrets value .yaml and support per path secret Detailed documentation on this can be found the [Vault Configuration](#) docs

Release 0.5.1

- Added back in the service name by default to the private cert generation pulled out of the previous release by accident. If the product was `pingaccess-admin` and release was `acme`, then the service name might be `acme-ping-access-admin`. This name by default will be added to the alternative hosts of the private certificate generation by default. Without this the pingaccess clustering will fail during setup.

Release 0.5.0

- [Issue #103](#) - Provide ability to add additional alt-names/alt-ips to private cert generation

Allow for a privateCert structure to contain optional arrays `additionalHosts` and `additionalIPs`:

```
pingaccess-admin:
  privateCert:
    generate: true
    additionalHosts:
    - pingaccess-admin.west-cluster.example.com
    - pa-admin.west-cluster.example.com
    additionalIPs:
    - 123.45.67.8
```

In addition, if the ingress for the product is enabled, the host(s) created for that ingress will also be added to the alt-names.

The above example (with an ingress) will create a cert used by pingaccess-admin containing:

```
Certificate:  
  Data:  
    ...  
    Signature Algorithm: sha256WithRSAEncryption  
    Issuer: CN=pingaccess-admin  
    ...  
    X509v3 extensions:  
      ...  
      X509v3 Subject Alternative Name:  
        DNS:rel050-pa-pingaccess-admin.ping-devops.com. pingaccess-admin.west-cluster.example.com,  
        DNS:pa-admin.west-cluster.example.com, IP Address:123.45.67.8
```

Release 0.4.9

- [Issue #104](#) — Update default global image tag to 2102 (Feb 2021)

Update the default global image tag in base values.yaml and remove edge from example yaml.

Release 0.4.8

- [Issue #100](#) — Change pingfederate-engine HPA to a default of disabled

Changing the default value `pingfederate-engine.clustering.autoscaling.enabled=false`, since the default CPU Request is set to 0.

Release 0.4.7

- [Issue #95](#) — Unable to set numerous Vault configuration options

Updated ability to add any hashicorp.vault annotation to the workload. As part of this effort, the existing name/values have been deprecated, however will continue to work for a period of time.

Updated details can be found in the Vault Config docs.

- [Issue #97](#) — Add the ability to add annotations to all resources generated similar to current support for Labels. This will allow deployers to specify additional annotations at either the global and/or product level. An example of the values.yaml would look like:

```
global:  
  annotations:  
    app.ping-devops.com/test: test-name  
  
pingaccess-admin:  
  annotations:  
    app.pingaccess/version: v1234
```

Additional cleanup of Notes.txt outputting detail of deployment.

Release 0.4.6

Minor follow-up update to cpu/memory request/limit sizes for init containers.

Release 0.4.5

- [Issue #89](#) — Update default workload resource cpu/memory request sizes.

Updating defaults to create a usage better reflecting actual memory usage by product. And minimizing amount of CPU needed as testing generally utilizes very little. Of course, it is definitely recommended that production deployments specify amount of cpu and memory required and limited to.

Current defaults are set to:

```
#-----
# Ping DevOps
#
# Description: All Ping Identity product images with integration
#-----
#
#          Product      Workload    cpu-R  cpu-L  mem-R  mem-L  Ing
#  -----
#  ✓ pingaccess-admin   deployment   0     2     1Gi    4Gi   false
#  ✓ pingaccess-engine  deployment   0     2     1Gi    4Gi   false
#  ✓ pingdataconsole    deployment   0     2     .5Gi   2Gi   false
#  ✓ pingdatagovernance deployment   0     2     1.5Gi  4Gi   false
#  ✓ pingdatagovernancepap deployment   0     2     .75Gi  2Gi   false
#  ✓ pingdatasync       deployment   0     2     .75Gi  2Gi   false
#  ✓ pingdelegator      deployment   0     500m  32Mi  64Mi  false
#  ✓ pingdirectory      statefulset 0     2     2Gi   8Gi   false
#  ✓ pingfederate-admin deployment   0     2     1Gi   4Gi   false
#  ✓ pingfederate-engine deployment   0     2     1Gi   4Gi   false
#
#  ✓ ldap-sdk-tools     deployment   0     0     0     0     false
#  ✓ pd-replication-timing deployment  0     0     0     0     false
#
#-----
```

Release 0.4.4

- [Issue #80](#) — Add support for importing a secret containing license into the container. Adds ability to add secret and configMap data to a container via a VolumeMount. A good use of this practice - bringing product licenses into the container.

Example of creating 3 volume mounts in container from secret and configMap

```
pingfederate-admin
  secretVolumes:
    pingfederate-license:
      items:
        license: /opt/in/instance/server/default/conf/pingfederate.lic
        hello: /opt/in/instance/server/default/hello.txt

  configMapVolumes:
    pingfederate-props:
      items:
        pf-props: /opt/in/etc/pingfederate.properties
```

In this case, a secret (called pingfederate-license) and configMap (called pingfederate-props) will bring in a couple of key values (license, hello) and (pf-props) into the container as specific files. The results will look like:

Example of kubectl describe of pingfederate-admin container

```
Containers:
  pingfederate-admin:
    Mounts:
      /opt/in/etc/pingfederate.properties from pingfederate-props (ro,path="pingfederate.properties")
      /opt/in/instance/server/default/conf/pingfederate.lic from pingfederate-license
      (ro,path="pingfederate.lic")
      /opt/in/instance/server/default/hello.txt from pingfederate-license (ro,path="hello.txt")
    Volumes:
      pingfederate-license:
        Type:      Secret (a volume populated by a Secret)
        SecretName: pingfederate-license
        Optional:  false
      pingfederate-props:
        Type:      ConfigMap (a volume populated by a ConfigMap)
        Name:     pingfederate-props
        Optional: false
```

Release 0.4.3

- [Issue #83](#) — Remove old pingdirectory tag check when creating service-cluster. This caused issues when creating a pingdirectory deployment with most recent tags (tags other than edge or 2012).

Release 0.4.2

- [Issue #79](#) — Adding support for product PingDataGovernance PAP
- [Issue #78](#) — Adding support to provide affinity definition to the workload of a product.

Example values.yaml to add podAntiAffinity to pingdirectory

```

pingdirectory:
  container:
    affinity:
      podAntiAffinity:
        # Add a hard requirement for each PD pod to be deployed to a different node
        requiredDuringSchedulingIgnoredDuringExecution:
          - labelSelector:
              matchExpressions:
                - key: app.kubernetes.io/name
                  operator: In
                  values:
                    - pingdirectory
            topologyKey: "kubernetes.io/hostname"
        # Add a soft requirement for each PD pod to be deployed to a different AZ
        preferredDuringSchedulingIgnoredDuringExecution:
          - weight: 1
            podAffinityTerm:
              labelSelector:
                matchExpressions:
                  - key: app.kubernetes.io/name
                    operator: In
                    values:
                      - pingdirectory
            topologyKey: "failure-domain.beta.kubernetes.io/zone"

```

Release 0.4.1

- Change default image tag to `2101` (January 2021).
 - Create private certs and keystore for use by images, only if the value `{product-name}.privateCert.generate=true`. Defaults are false.
 - Helm will generate the a `tls.crt` and `tls.key`, place it into a kubernetes secret called `{release-productname}-private-cert`.
 - Mount the secret into the image under `/run/secrets/private-cert`
 - An init container will pull the `tls.crt` and `tls.key` into a pkcs12 keystore and place it into a file `/run/secrets/private-keystore/keystore.env` that will be mounted into the running container.
 - When the container's hooks are running, it will source the environment variables in this `keystore.env`. The default variables set are:
 - `PRIVATE_KEYSTORE_PIN={base64 random pin}`
 - `PRIVATE_KEYSTORE_TYPE=pkcs12`
 - `PRIVATE_KEYSTORE={pkcs12 keystore}`
- yaml to generate a private cert/keystore for pingaccess-admin:
- ```

pingaccess-admin:
 privateCert:
 generate: true

```
- Example of created `/run/secrets/private-keystore/keystore.env`

```
PRIVATE_KEYSTORE_PIN=nrZmV4XdfK....
PRIVATE_KEYSTORE_TYPE=pkcs12
PRIVATE_KEYSTORE=MIIJgQIBAzCCUcGC....
```

- Added support for PingAccess clustering between pingaccess-admin and multiple pingaccess-engine containers.
  - See everything.yaml for example of deploying a PingAccess cluster using PingFederate/PingDirectory to authenticate
  - It is required to either:
    - generate the private cert (see above) with the value of pingaccess-admin.privateCert.generate=true or
    - provide your own cert secret called {release-productname}-private-cert containing a valid tls.crt and tls.key.
  - Enable both the pingaccess-admin and pingaccess-engine helm chart products
    - Example values to create a clustered pingaccess:

```
pingaccess-admin:
 enabled: true
 privateCert:
 generate: true
 envs:
 SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git
 SERVER_PROFILE_PATH: baseline/pingaccess
pingaccess-engine:
 enabled: true
 envs:
 SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git
 SERVER_PROFILE_PATH: baseline/pingaccess
pingfederate-admin:
 enabled: true
 envs:
 SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git
 SERVER_PROFILE_PATH: baseline/pingfederate
 container:
 waitFor:
 pingdirectory:
 service: ldaps
pingfederate-engine:
 enabled: true
 envs:
 SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git
 SERVER_PROFILE_PATH: baseline/pingfederate
pingdirectory:
 enabled: true
 envs:
 SERVER_PROFILE_URL: https://github.com/pingidentity/pingidentity-server-profiles.git
 SERVER_PROFILE_PATH: baseline/pingdirectory
```

## Release 0.4.0

- Support availability of PingDirectory pods through the cluster headless kubernetes service. Allows for PingDirectory nodes to find one another during the replication enable/init process.

Adds following to pingdirectory-cluster:

```
metadata:
 annotations:
 service.alpha.kubernetes.io/tolerate-unready-endpoints: "true"
spec:
 publishNotReadyAddresses: true
```

## Release 0.3.9

- Fixed the default wait-for service name on pingfederate-engine (admin → https).
- Changed default on readiness command to check for readiness every 5 seconds rather than 30. This allows for availability on some services, such as PingFederate which is normally ready in 30 sec.

## Release 0.3.8

- [Issue #56](#) — Improved Default Naming on Global vars - PORTs
- [Issue #56](#) — Improved Default Naming on Global vars - PORTs
- In Release 0.3.6, global-env-vars were created for PORTs. The naming structure used was complex and difficult, primarily because a product can have several ports open on a particular private and public host. The format will be more consistent as defined by the following:

```
{product-short-code with type}_{public or private}_{hostname or port}{_service if port}
```

An example with PD might look like (note the service names of `https` and `data-api`):

```
PD_ENGINE_PUBLIC_PORT_HTTPS: 443
PD_ENGINE_PUBLIC_PORT_DATA_API: 1443

PD_ENGINE_PRIVATE_PORT_HTTPS: 443
PD_ENGINE_PRIVATE_PORT_DATA_API: 8443
```

- [Issue #62](#) — When creating configMapRef's, take into account the proper release name to include ConfigMapRef names in workloads were not consistent with the ConfigMaps created by default when taking into account the `addReleaseNameToResource` setting of prepend, append or none. This fixes that issue ensuring that config maps are consistent.
- Added global-env private/public host/port for PingDataConsole, which was missing.
- Changed the default pingfederate-admin `admin` service name to `https` to reduce confusion.
- Changed the default pingfederate-engine `engine` service name to `https` to reduce confusion.

## Release 0.3.7

- Fixes issue with service -vs- ingress name on creation of ingress to service mapping. Resolves issue #57.

## Release 0.3.6

- Cleaning up and making services/ingresses easier to use together. Incorporating all the ports used in both a service and ingress into the same location of the [Service Configuration](#).

The example below shows a container/service/ingress and how to specify the ports at each level.

- `containerPort` → Replaces `targetPort`
- `servicePort` → Replaces `port`
- `ingressPort` → New entry

```
services:
 api:
 containerPort: 8443 <--- changed from targetPort
 servicePort: 1443 <--- changed from port
 ingressPort: 443 <--- new. moved from ingress
 dataService: true
 data-api:
 containerPort: 9443 <--- changed from targetPort
 servicePort: 2443 <--- changed from port
 ingressPort: 2443 <--- new. moved from ingress
 dataService: true
 ingress:
 hosts:
 - host: pingdirectory.example.com
 paths:
 - path: /api
 backend:
 serviceName: api <--- changed from servicePort
 - path: /directory/v1
 backend:
 serviceName: data-api <--- changed from servicePort
```

Additionally, `global-env-vars` will be created for each of these ports. If the name of the product is `PROD`, the the following ports would be created:

```
PROD_API_PRIVATE_PORT="1443" # This is the servicePort
PROD_API_PUBLIC_PORT="443" # This is the ingressPort
PROD_DATA_API_PRIVATE_PORT="2443"
PROD_DATA_API_PUBLIC_PORT="2443"
```

- Fixed missing `USER_BASE_DN` setting in simple-sync.yaml example.

## Release 0.3.5

- Allowing config values to determine use of init containers to wait-for other chart products. For each product, you can now provide a `waitFor` structure providing the name and service that should be waited on before the running container can continue. This will basically inject an initContainer using the PingToolkit wait-for utility until it can `nc host:port` before continuing.

PingFederate Admin waiting on pingdirectory ldaps service to be available:

```
pingfederate-admin:
 container:
 waitFor:
 pingdirectory:
 service: ldaps
 pingdatagovernance:
 service: https
```

- By default, the `pingfederate-engine` will `waitFor` `pingfederate-admin` before it starts.

## Release 0.3.4

- Adding init container to PingFederate Admin to wait-for PingDirectory's LDAPs port if the `pingdirectory.enabled=true`. This fixes an issue that keeps PingFederate Admin from starting when it's dependent on PingDirectory. In the case that PingFederate isn't dependent on PingDirectory and it is still enabled, it will simply delay the start time of PingFederate admin. A future version will allow for specifying a list of services to wait-for so this can be turned off/on by deployer.
- Moved the `securityContext` settings added to release 0.3.3 from the container to the workload, as that is the proper place to use them. Required for use of `fsGroup` setting.

## Release 0.3.3

- Adding the ability for a deployer to add a `securityContext` to the containers. Currently, there are warning messages in the images when an outside-in pattern is used (i.e. `securityContext` is set). Also, many of the default ports require privileged access, so care should be taken along with testing to ensure the containers start up fine. Additional, one should not change the security context when doing an upgrade or using a PCV from a previous deployment.

An example `securityContext` that can be used might look like:

```
global:
 container:
 securityContext:
 allowPrivilegeEscalation: false
 capabilities:
 drop:
 - ALL
 runAsGroup: 1000
 runAsNonRoot: true
 runAsUser: 100
```

By default, the `values.yaml` in the chart will set the `securityContext` to empty:

```
global:
 container:
 securityContext: {}
```

## Release 0.3.2

- Replaced init container on pingfederate-engine to use pingtoolkit rather than 3rd party curlimage. Additionally added resource constraints and security context to this init container.
- Removed hardcoded SERVER\_PROFILE\_BRANCH set to master, relying on git repo default branch
- Cleaned up pingdelegator values. public hostnames for pingfederate and pingdirectory built based off of ingress hostnames, part of `{release-name}-global-env-vars` configmap.
- Removed default nginx annotations of ingress resources. If an nginx controller is used for ingress, the following ingress annotations should be included:

### Warning

By removing the following annotations from the default, use of current config values will result in no ingress being set. You must add these in via your .yaml file or via separate --set settings.

```
global:
 ingress:
 annotations:
 nginx.ingress.kubernetes.io/backend-protocol: "HTTPS"
 kubernetes.io/ingress.class: "nginx-public"
```

## Release 0.3.1

- Added container envFrom for `{release-name}-env-vars` back as optional. Fixes breaking change from 0.2.8 to 0.2.9 for those that used this configmap.
- Added ability for deployer to add their own envFrom's via their values.yaml. An example (adding an optional configmap/ secrets to all products). Just change global to the name of the product to only have that product use the references.

```
global:
 container:
 envFrom:
 - configMapRef:
 name: my-killer-configmap
 optional: true
 - secretRef:
 name: my-killer-secrets
 optional: true
```

## Release 0.3.0

- Consolidate deployment/stateful set templates to a single workload template.
- Changes to values.yaml:
  - Created a workload map under global (see below)
  - Moved old deployment information under workload

- Moved old statefulSet information under workload
- Updated `pingfederate-admin` to reflect new workload
- Updated `pingdirectory` to reflect new workload
- Allows for any product to be run as a deployment or statefulSet

## Warning

Using `workload.type=StatefulSet` will create `pvc` resources and allow for persistence on restarts of containers. This is helpful during development. Be aware that the `pvc` resources will need to be deleted to startup a fresh copy of the product images.

```
global: workload: type: Deployment # Can be Deployment or StatefulSet (see warning above)

deployment:
 strategy:
 type: RollingUpdate # Can be RollingUpdate or Recreate
 rollingUpdate:
 maxSurge: 1
 maxUnavailable: 0

 statefulSet:
 partition: 0 # Used for canary testing if n>0

 persistentvolume:
 enabled: true
 #####
 # For every volume defined in the volumes list, 3 items will be
 # created in the StatefulSet
 # 1. container.volumeMounts - name and mountPath
 # 2. template.spec.volume - name and persistentVolumeClaim.claimName
 # 3. spec.volumeClaimTemplates - persistentVolumeClaim
 #
 # https://kubernetes.io/docs/concepts/storage/persistent-volumes/
 #####
 volumes:
 out-dir:
 mountPath: /opt/out
 persistentVolumeClaim:
 accessModes:
 - ReadWriteOnce
 storageClassName:
 resources:
 requests:
 storage: 4Gi
```

- Renamed template files in pinglib from .yaml to .tpl
- Added `terminationGracePeriodSeconds` to container to support setting in values
- Added `serviceAccountName` to vault.hashicorp to specify to the container what service account can be used to authenticate to the Hashicorp Vault Injector